



# Process Abstraction

2023-24 COMP3230B

# Contents

- What is a “process”?
- Represent a “process”
  - Resources use by a process
  - Process states
- Important OS data structures
- Operations on processes

# Learning Outcome

- ILO 2a - **explain** how OS manages **processes**/threads and discuss the mechanisms and policies in efficiently sharing of CPU resources.

# Reading & Reference

- Required Reading
  - Operating Systems: Three Easy Pieces by Arpaci-Dusseau et. al
    - Chapter 4, Abstraction: The Process
      - <http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-intro.pdf>
    - Chapter 5, Interlude: Process API
      - <http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-api.pdf>

# Process vs. Program

- Program itself is a lifeless thing
- What is a Process?
  - A process is a program in execution
    - **An instance** of a program running in a computer
    - A process is an entity that can be assigned to and executed on a CPU
  - A unit of activity characterized by
    - the execution of a sequence of instructions,
    - with an execution state, and
    - an associated set of system resources

# Process – an abstraction

- To represent a running program, the OS needs to keep the following information:

The memory that the process can access (or reference) is part of the process

Memory

## Examples

Program code

Data

Resources in  
use

I/O in use

Physical memory

During execution, process updates/ stores data in CPU registers, e.g., program counter, stack pointer, etc.

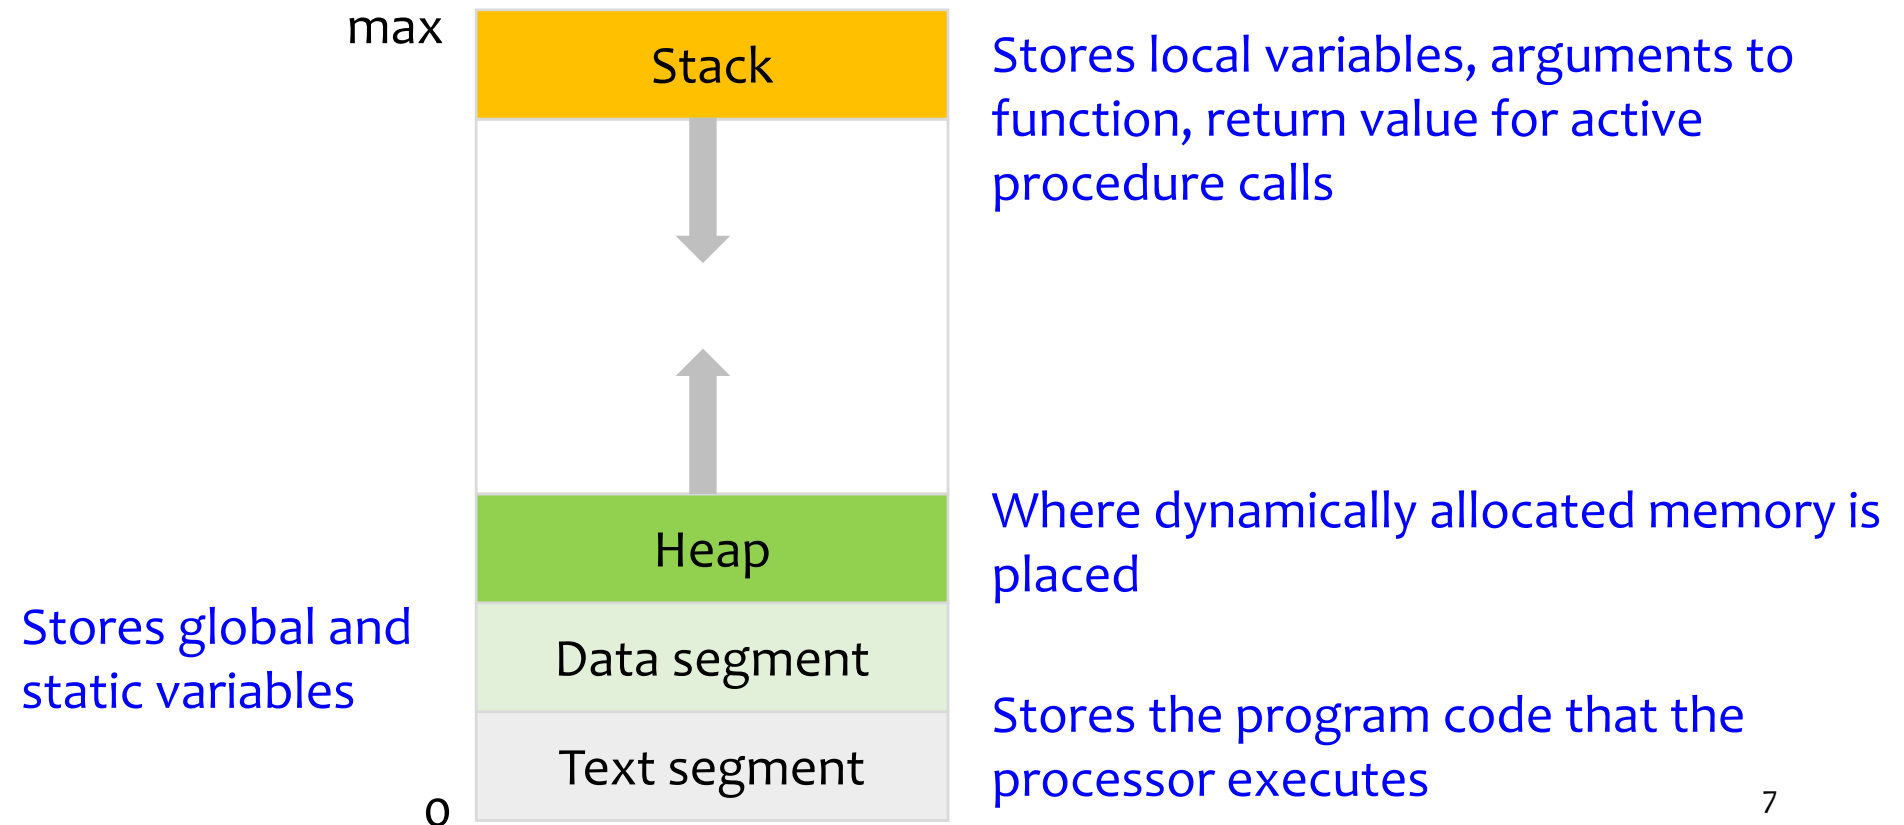
Execution  
state

Register set

Current process  
state

# Process – Address Space

- The process's view of its memory is called the **address space**, which is a range of memory locations (or a range of memory addresses)
- Process address space consisting of a few “regions”:

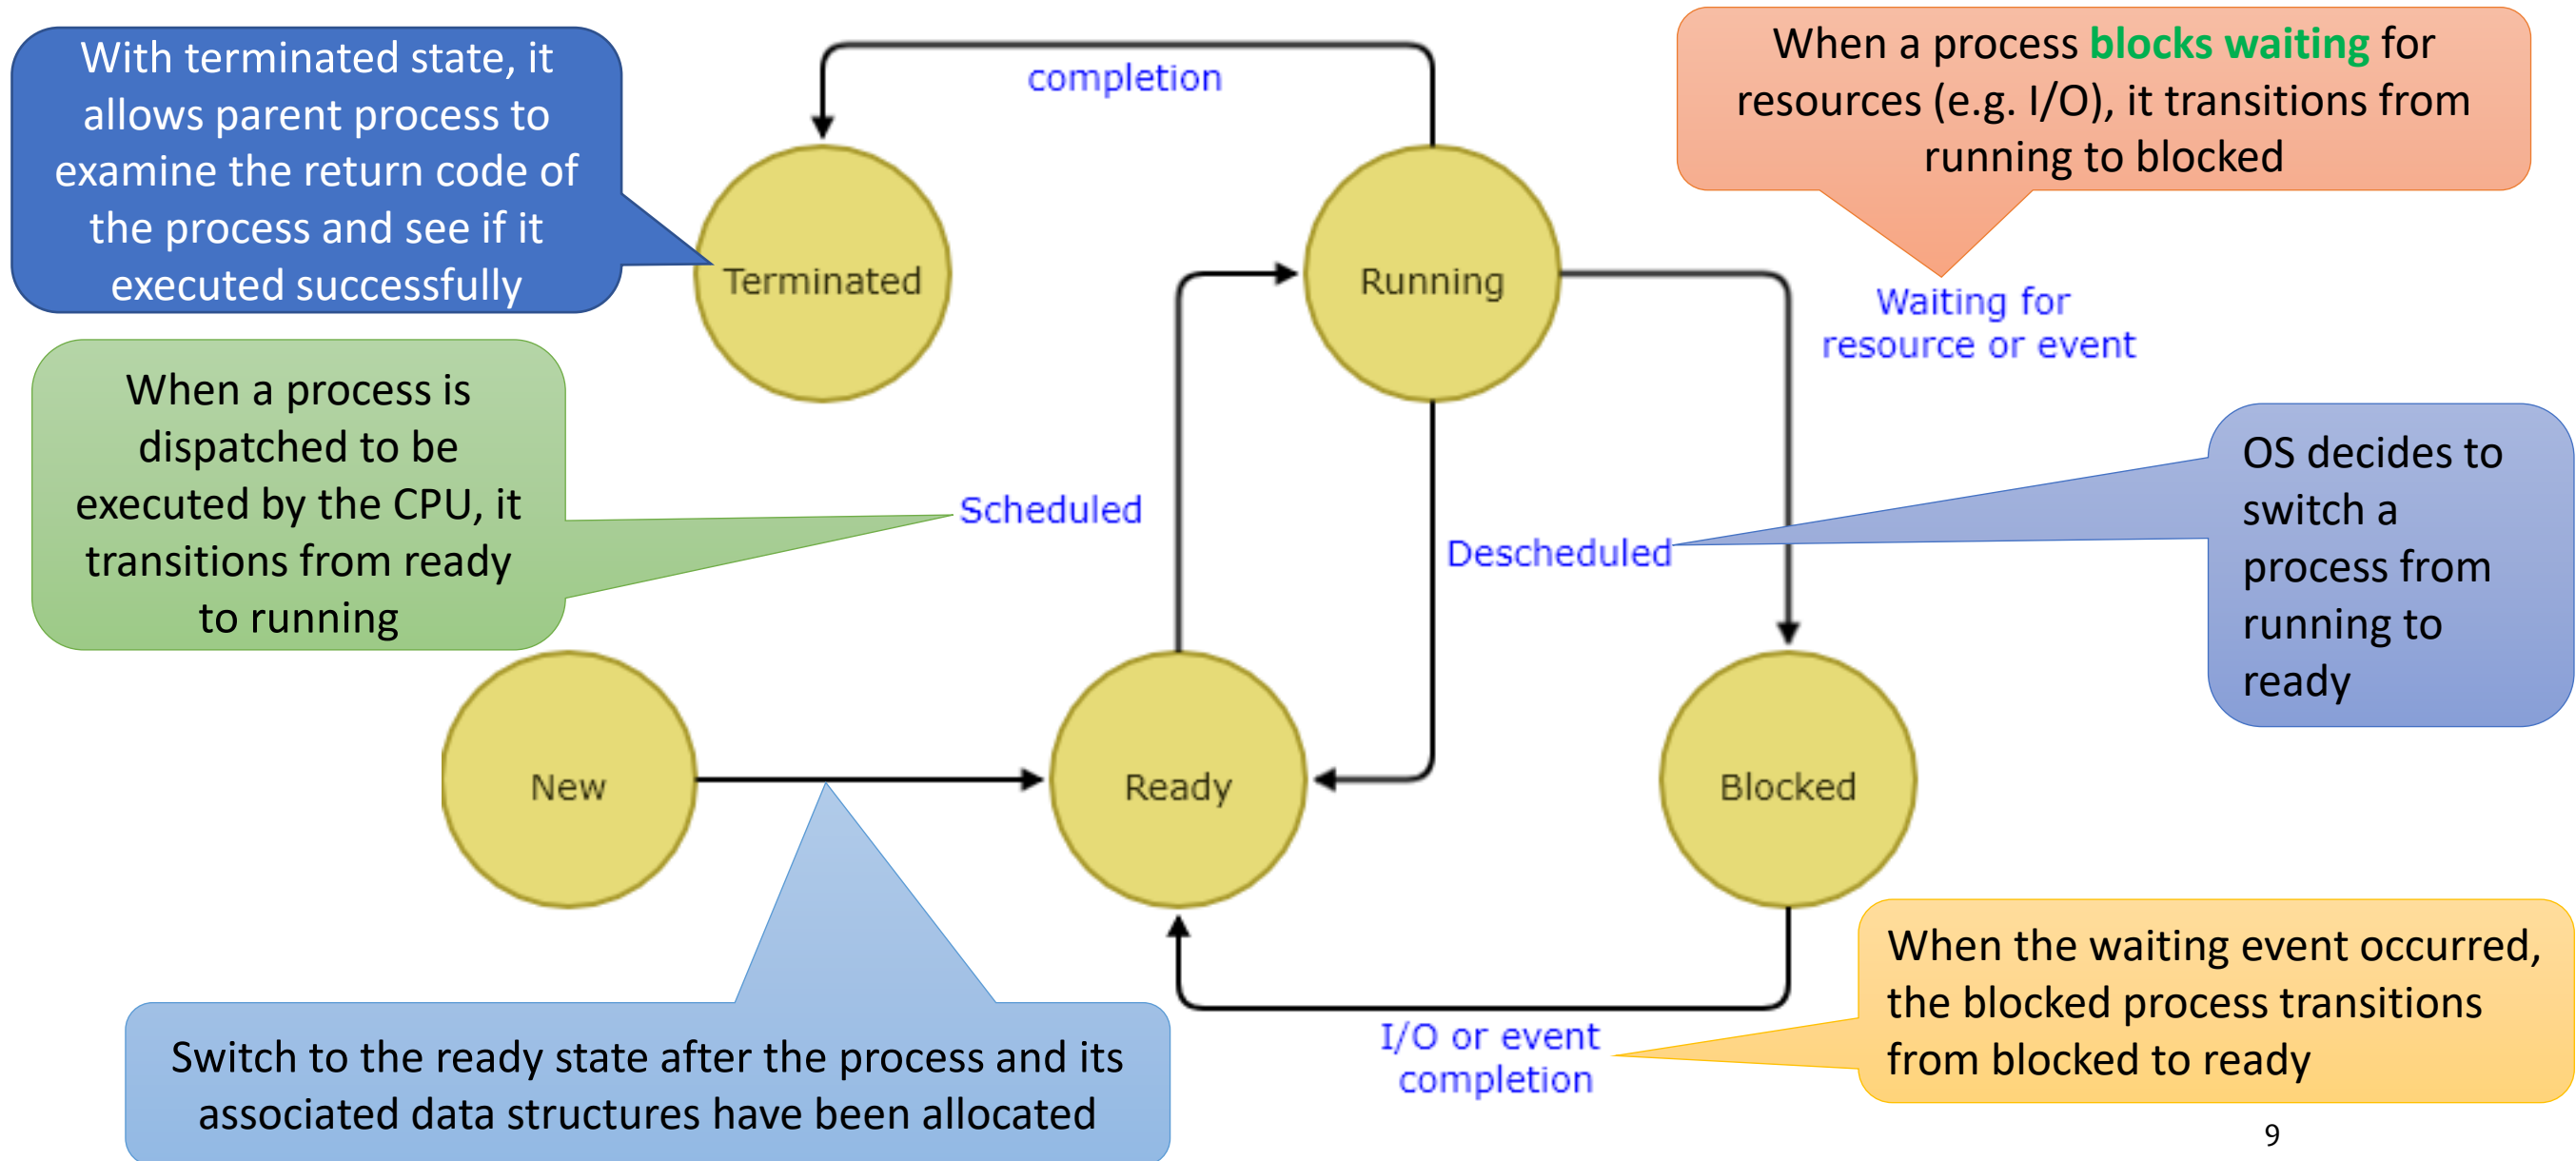


# Process – Process States

- In the process execution life cycle, it moves through a series of discrete process states.
- Process state - an indicator of the nature of the current activity of a process
  - **new (initial)**: The process is just being created
  - **running**: The process is executing on a processor
  - **blocked**: The process is **waiting for** some event (e.g., I/O or communication) to happen before it can proceed
  - **ready**: The process **is ready to run** on a processor and **is waiting** to be assigned to a processor
  - **terminated (final / zombie)**: The process has finished execution but has not yet been cleaned up; why not just discard it?



# Life Cycle of a Process



# Process Control Block

- To manage a process, OS makes use of a **data structure** to maintain information about a process
  - Process Control Block (**PCB**) or Process Descriptor
- PCB typically includes
  - Process identification number (PID) - a unique ID
  - The current process state
  - The program counter - indicates the address of **next** instruction
  - Register context - a snapshot of the register contents in which the process **was last running before** it **transitioned out of the running state**
  - Scheduling information - process priority, pointers to scheduling queues, etc.

## Process Control Block (2)

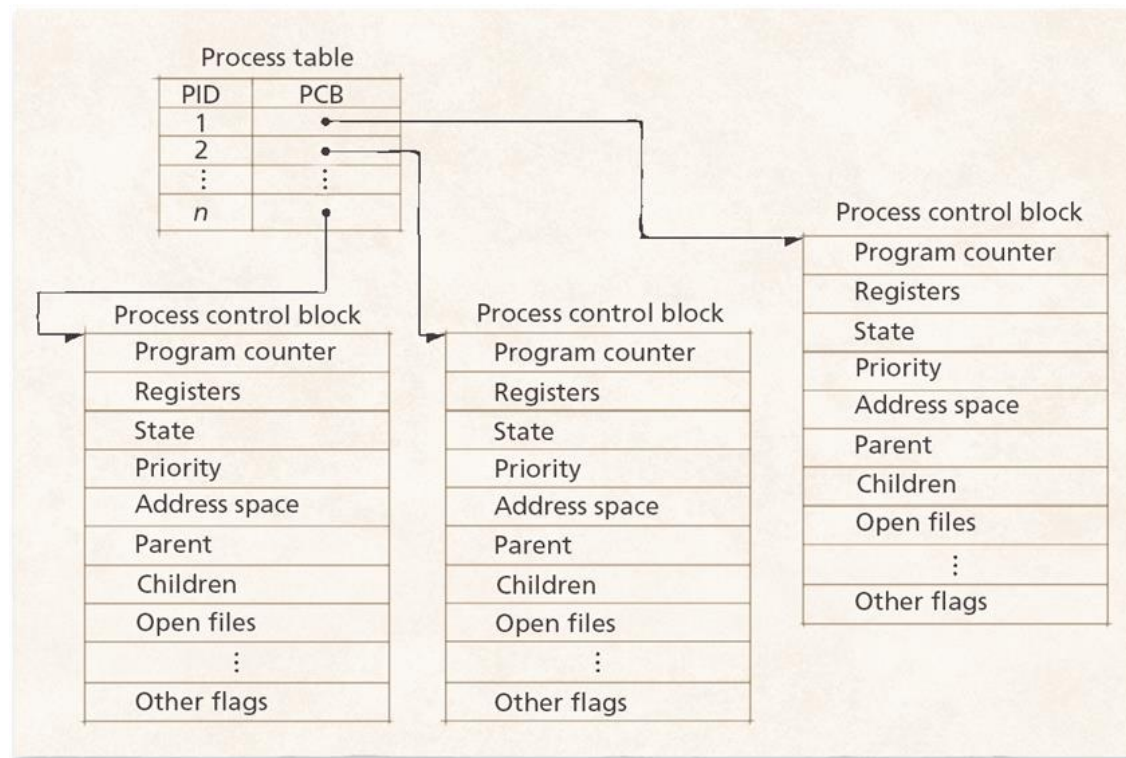
- Credentials - determines the resources this process can access
- Memory Management information - concerning memory areas allocated to the process
- Accounting information - CPU usage statistics, time limits, etc.
- A pointer to the process's parent process
- Pointers to the process's child processes
- Pointers to allocated resources

:

Example: Linux process descriptor  
***struct task\_struct***  
in /usr/src/linux/include/**linux/sched.h**  
at around 800 lines of statements

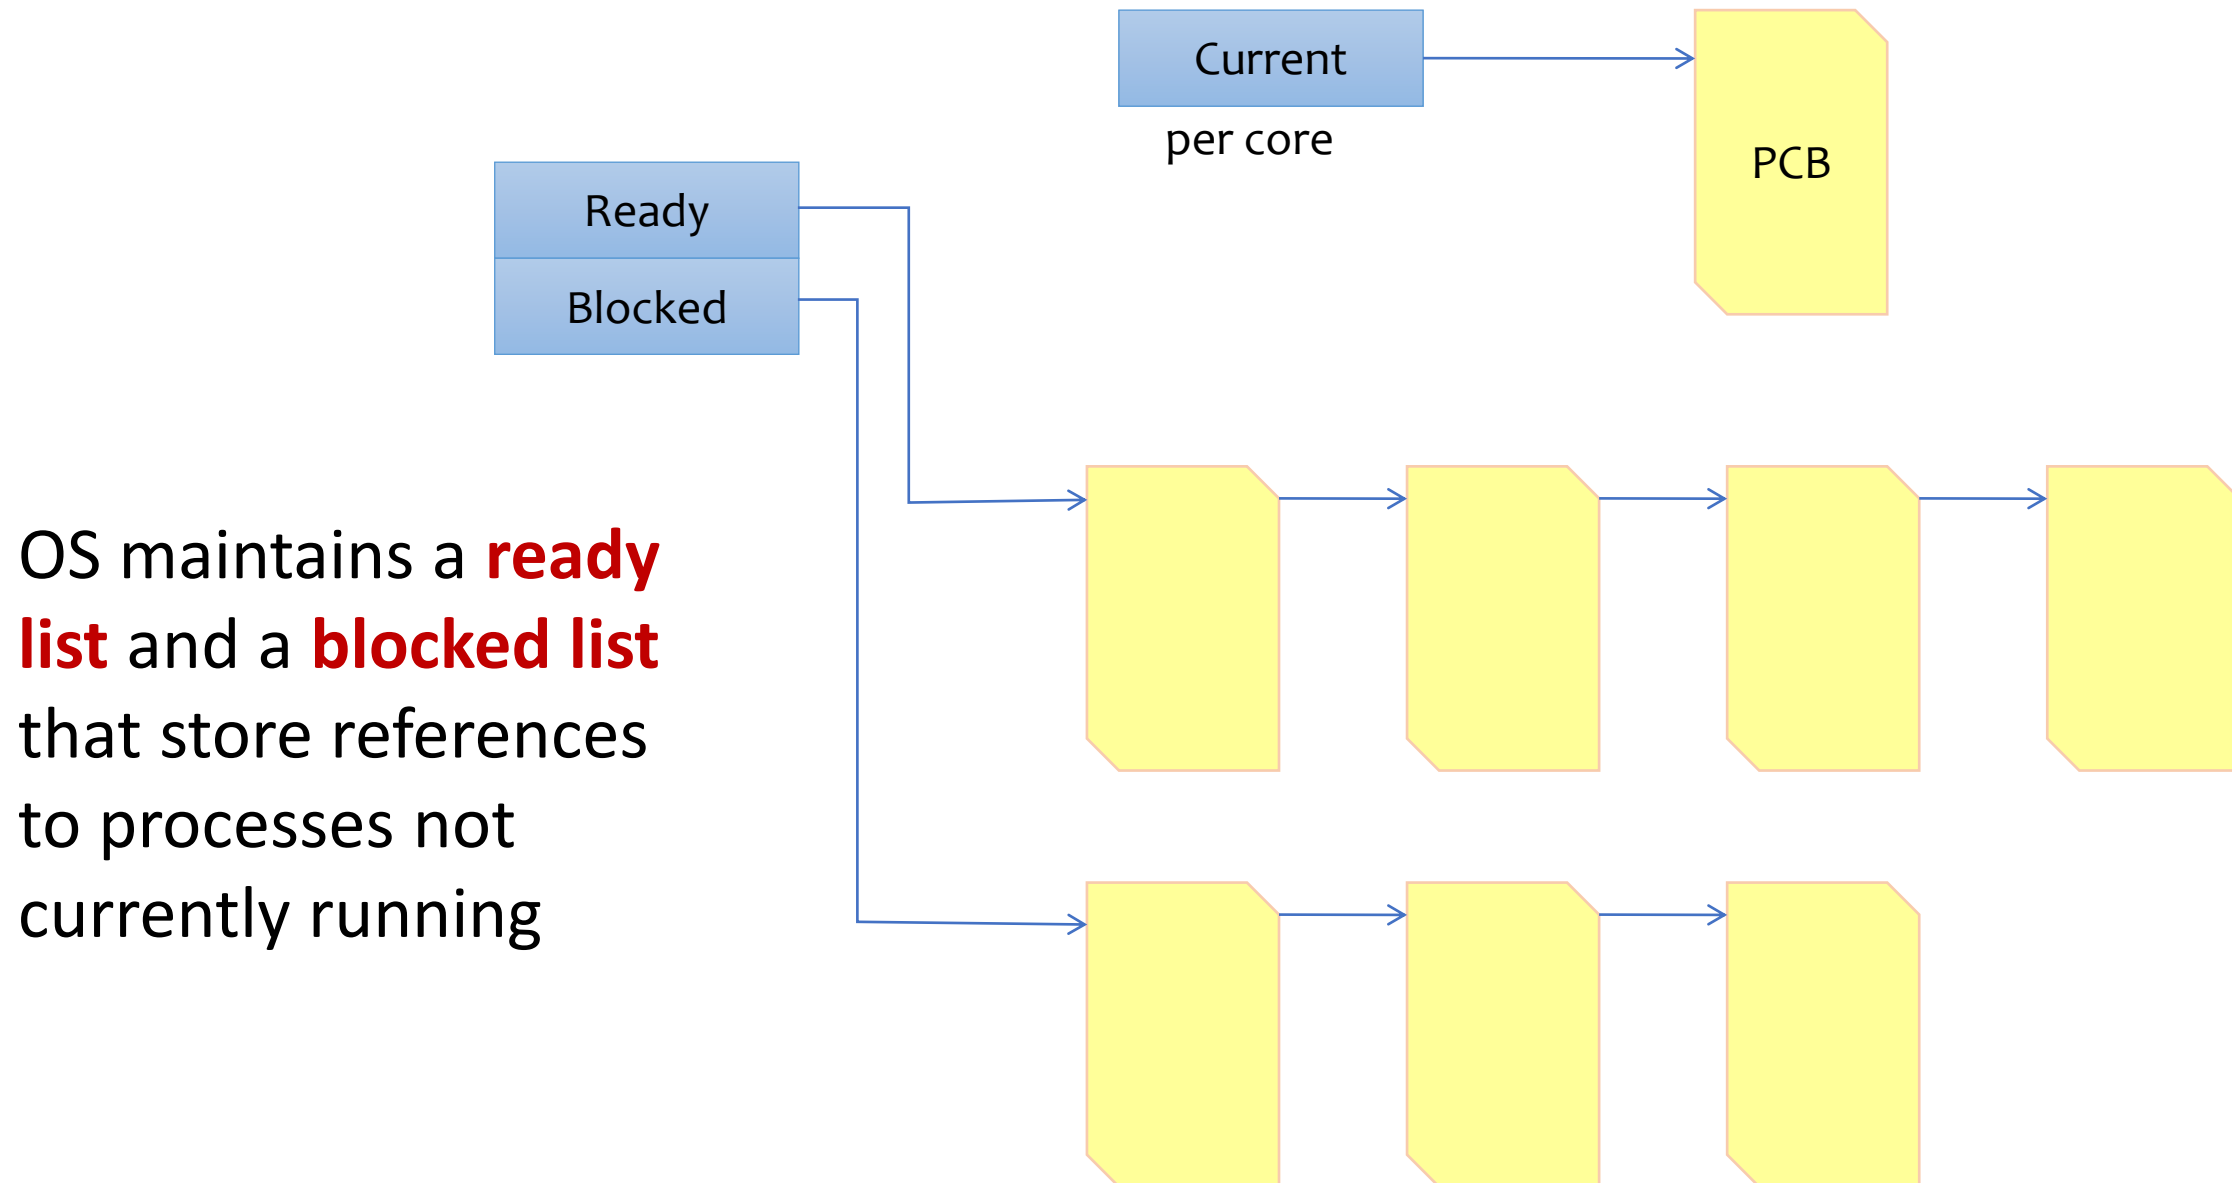
# Process Table

- To manage many processes, OS needs a way to **quickly access** a process's PCB



- OS keeps pointers to each process's PCB in a table
  - Example - Linux "process table" is organized in a form of hashed table
- When a process is "completely" terminated, OS removes the process from the process table and frees all of the process's resources

# Process List Structures



# ps - Show Processes information

- You can list the processes' details in Linux/Mac OS X by using the **ps** command
- Usage: ps [option]
  - When executed without any options, only processes that are associated with the current terminal are shown.
- Use **man-page** to learn how to use
- Some Useful options (Linux)
  - -e: Select all processes
  - -f: Show in full format
  - w: Wide output
  - f: ASCII-art process hierarchy (forest)

# pstree - Show the relations between processes

- A Linux system program similar to ps
  - Processes are organized in hierarchy
    - As every process has a parent process, their relationship can be viewed by using pstree
- Usage: pstree [option]
  - Use man-page to learn how to use
  - Example:

```
atctam@workbench2:~$ ps
  PID TTY          TIME CMD
 3120691 pts/1    00:00:00 bash
 3123097 pts/1    00:00:00 ps
atctam@workbench2:~$ pstree -s 3120691
systemd—sshd—sshd—sshd—bash—pstree
atctam@workbench2:~$ pstree -sp 3120691
systemd(1)—sshd(205)—sshd(3120553)—sshd(3120626)—bash(3120691)—pstree(3123102)
atctam@workbench2:~$
```

# Operations on Processes

- Operating systems provide fundamental services to processes including:
    - Creating processes
    - Destroying processes
    - Suspending processes
    - Resuming processes
    - Changing process's priority (for scheduling)
    - Waiting for a process (parent process waits for the child process)
    - Check process's status
    - Interprocess communication (IPC)
- :



# How to create a new process?

- A process spawns a new process
  - The process that creates a new process is now called the **parent** process
  - The newly created process is called the **child** process
    - The child can also create other processes, thus forming a tree of processes
- When the parent process is destroyed/terminated, modern operating system typically
  - Keeps the child processes and allows them to proceed independently of the terminated parent process

# Creating process

- Allocate memory for the process
  - Space for PCB must be allocated
- Initialize the process control block
  - Assign a unique process ID
  - Save the process ID, parent ID
  - Set program counter and stack pointer to appropriate values
- Create other data structures
  - Memory, files, accounting
- Set the process state to Ready and put it to the Ready queue

# How to Create a Process

- Unix
  - fork() system function
    - A system call that **creates** a new process by **duplicating** the calling process
  - exec () family of functions
    - OS replaces the current program image with a new program image
- Windows API
  - CreateProcess() function
    - Creates a new process and its primary thread and loads program for execution
    - [http://msdn.microsoft.com/en-us/library/ms682512\(VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ms682512(VS.85).aspx)
  - CreateProcess() is similar to fork() + exec ()

# Process Termination

- Process executes last statement and asks the operating system to delete itself
  - `exit()` or `return` from `main()`
- A process may terminate **involuntarily**
  - Parent may terminate execution of children processes (by sending a termination **signal**)
  - A number of **error and fault** conditions can lead to termination of process
- Return termination status from child to parent
  - Parent can obtain this info by calling `wait()`, `wait4()`, `waitid()` or `waitpid()`
  - The terminated process' resources are **de-allocated** by OS **afterward**

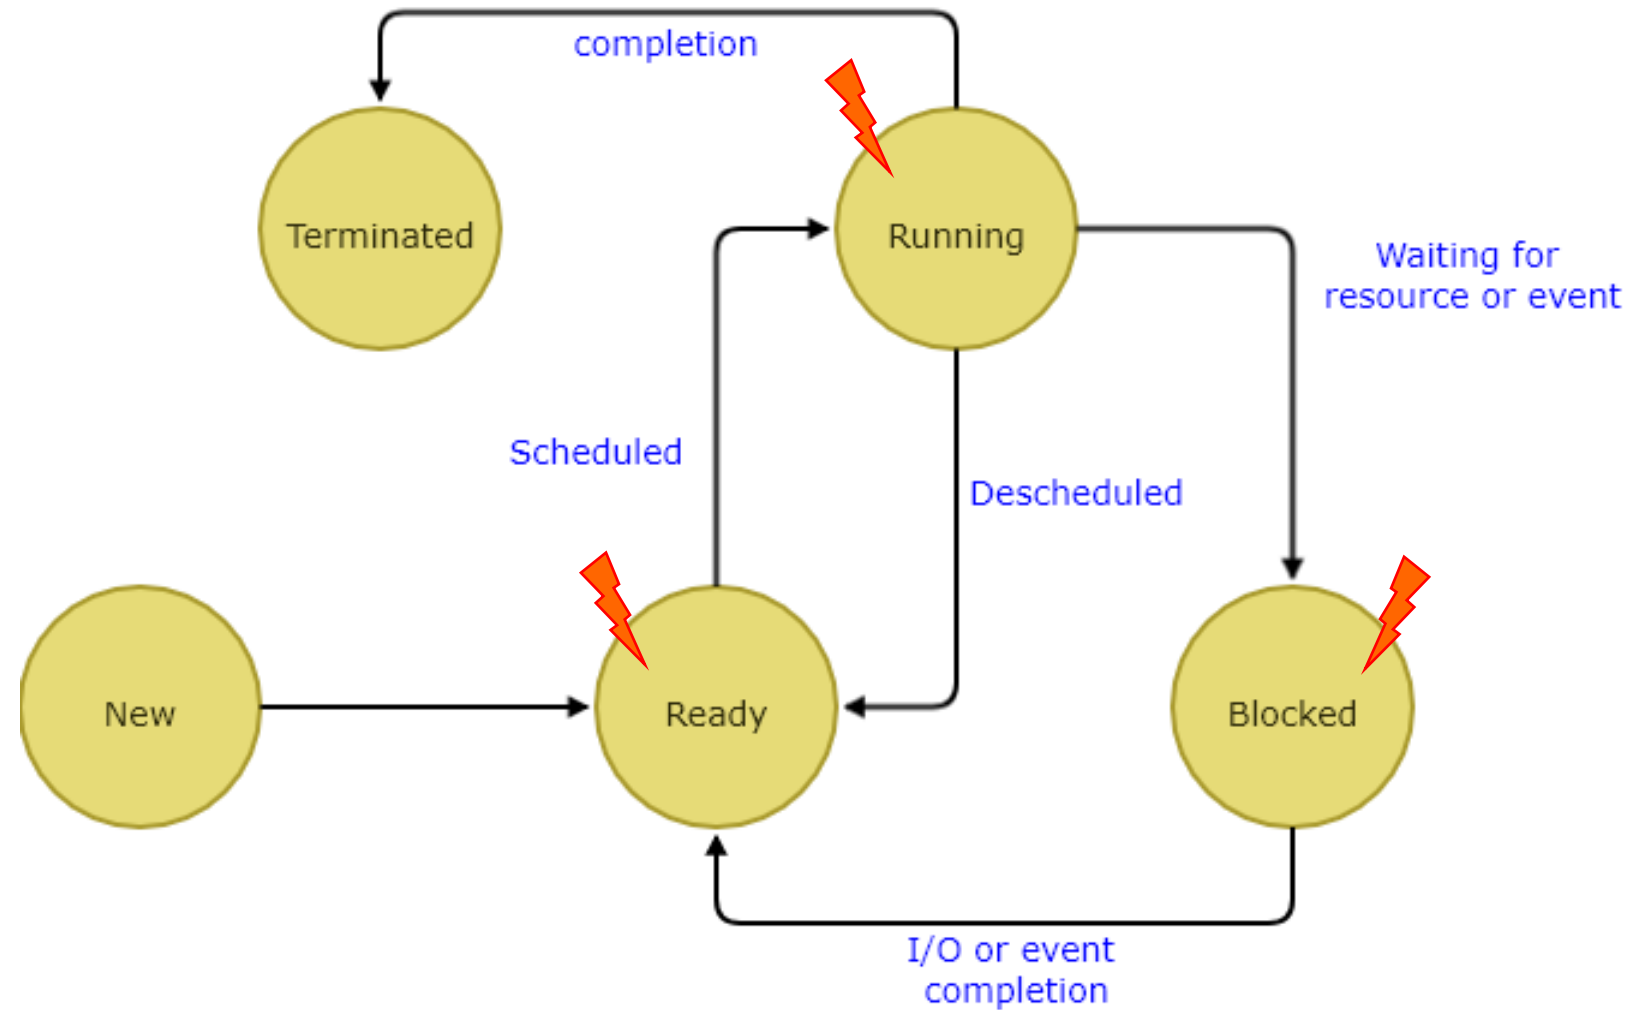
# Zombie Process (UNIX)

- When a process is in the terminated state, we call it a “zombie” process
  - When a process exits, OS still **keeps** the PCB of the process, so that the **parent can get** the information (if necessary)
    - exit status
    - resource usage
  - Parent uses waitpid() (and others) to inform OS to completely remove the zombie child process

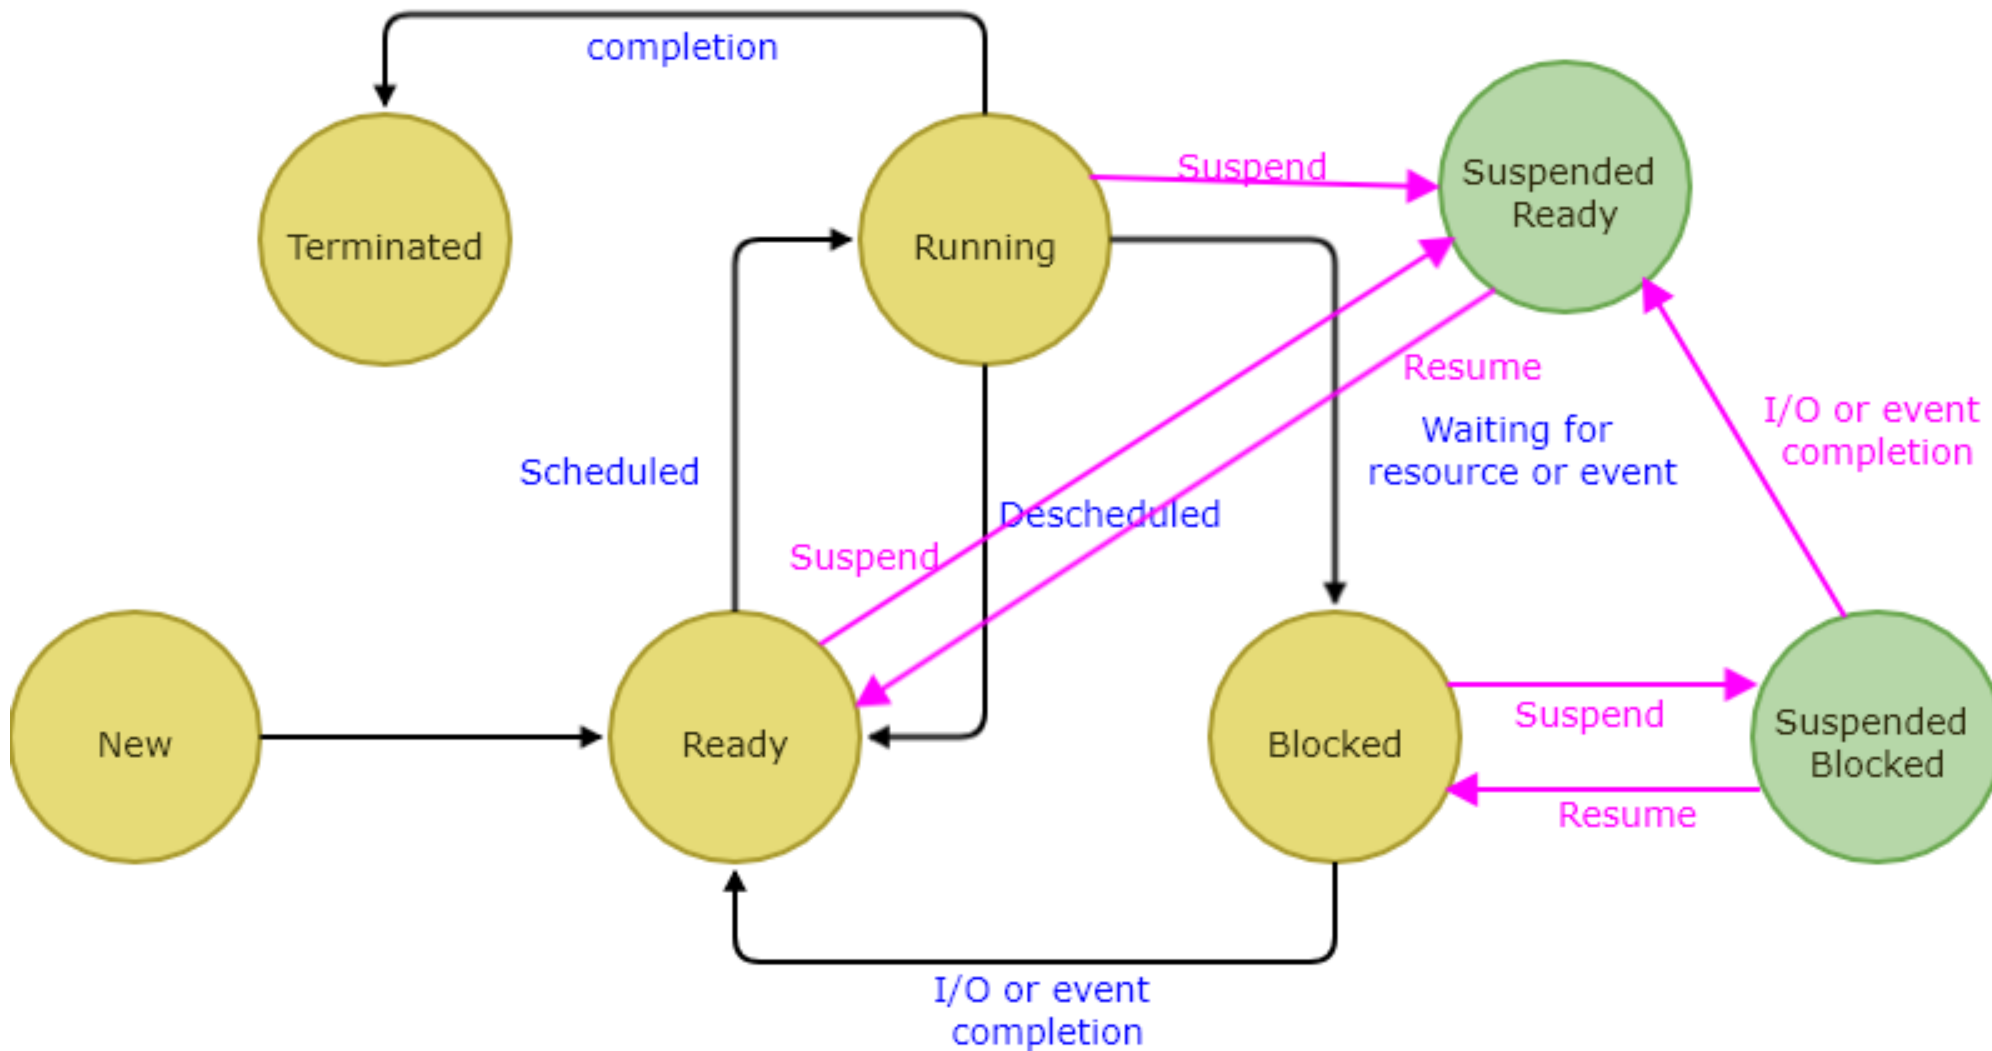
# Suspending a Process

- Temporarily deactivate the process, such that it is not being considered for processor scheduling
- Why doing so:
  - Upon user request
  - Request by the parent process
  - OS may decide to suspend a **blocked** process so as to free up the memory occupied by that process for another ready process
- A suspended process **must be** resumed by another process
- Difference between suspension and blocked
  - **blocking** is triggered by **internal activity** of the process, while **suspend** is coming from **external**

# Suspending a Process



# Suspending a Process



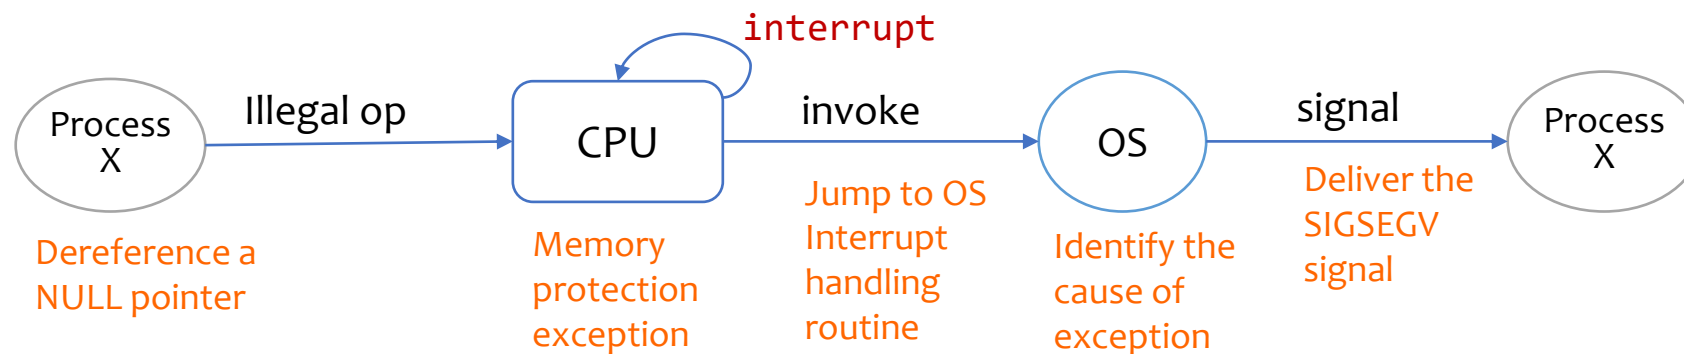


# Signals

- Signals are used in UNIX systems to **notify** a process that an **event** has occurred
  - Sometimes being referred as “**Software Interrupt**”
- Basically implemented as system calls such as kill(), signal(), sigaction(), raise(), pause(), sigsuspend(), etc.
- Each signal is represented by a value/symbolic name
  - SIGINT: value = 2, generated when Ctrl-c is pressed
  - SIGCHLD: value = 17, generated when child process finishes execution or is terminated


## Signals (2)

- A signal is generated by one **software entity** (in the occurrence of an event) to a **target software entity**
  - **Synchronous signal** – is triggered by the **current instruction** of the current running process itself and **is delivered** to **that** process by the OS immediately
    - e.g., illegal memory access, division by zero



# Signals (3)

- **Asynchronous signals** – are generated by external events / activities, which are **not triggered by the current activity** / action of the target process at the time of receiving the signal
  - i.e., arrive at unpredictable times during execution of the program
  - e.g., by the timer alarm
  - e.g., parent process using kill() system call to kill the child process



# Signals (4)

- A process can decide whether it wants to **catch**, **ignore** or **mask** a signal
  - **Catching** a signal involves specifying a routine (**signal handler**) in advance so that the OS will invoke that handler when the process receives that signal
    - the **signal()** or **sigaction()** system calls can be used by the program to specify the signal handler routine to the OS
  - **Catching** – Using OS's default action to handle the signal
  - **Ignore** – Inform OS that it does not want to handle that signal
  - **Masking** a signal is to instruct the OS not to deliver signals of that type until the process clears the signal mask
- SIGKILL and SIGSTOP cannot be caught, blocked or ignored

# Signals (5)

- How OS determines what a process will respond to a particular signal
  - A process's PCB contains a pointer to **a vector of signal handlers** (logically order by the signal number)
  - Each entry corresponds to the handler function for that entry (signal)
- A child process inherits the setting from its parent
- However, if use `exec...()` function to load a new program image, any signals that have the custom-made handlers **will be reset** to default setting

# Summary

- To manage and control a running process (application), OS needs some mechanisms to keep track on the current status of the process
  - Various data structures are needed
    - Process control block (PCB)
    - Process table and lists
- OS provides a set of operations for us to work with processes
- Signals were introduced in Unix systems to allow interactions between User Mode processes; the kernel also uses them to notify processes of system events.