


Virtualizing the CPU

2023-24 COMP3230B



Contents

- System Calls
- Process switch
- Interrupt

Related Learning Outcome

- ILO 2a - explain how OS manages processes/threads and discuss the mechanisms and policies in efficiently sharing of CPU resources.

Reading & Reference

- Required Reading
 - Chapter 6, Mechanism: Limited Direct Execution, Operating Systems: Three Easy Pieces by Arpaci-Dusseau et. Al
 - <http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-mechanisms.pdf>

Process Control


- How to perform restricted operations?
 - **Mechanism:** System Calls
- How to provide illusion of having many CPUs?
 - Virtualizing the CPU
 - **Mechanism:** Context Switching
- How to regain control of the CPU?
 - Voluntary release of CPU
Mechanism: System Calls
 - Involuntary release of CPU
Mechanism: Interrupt

Transparent:

- Process does not know when it is running and when it is not
- Programmer does not need to worry about this situation

To request OS services

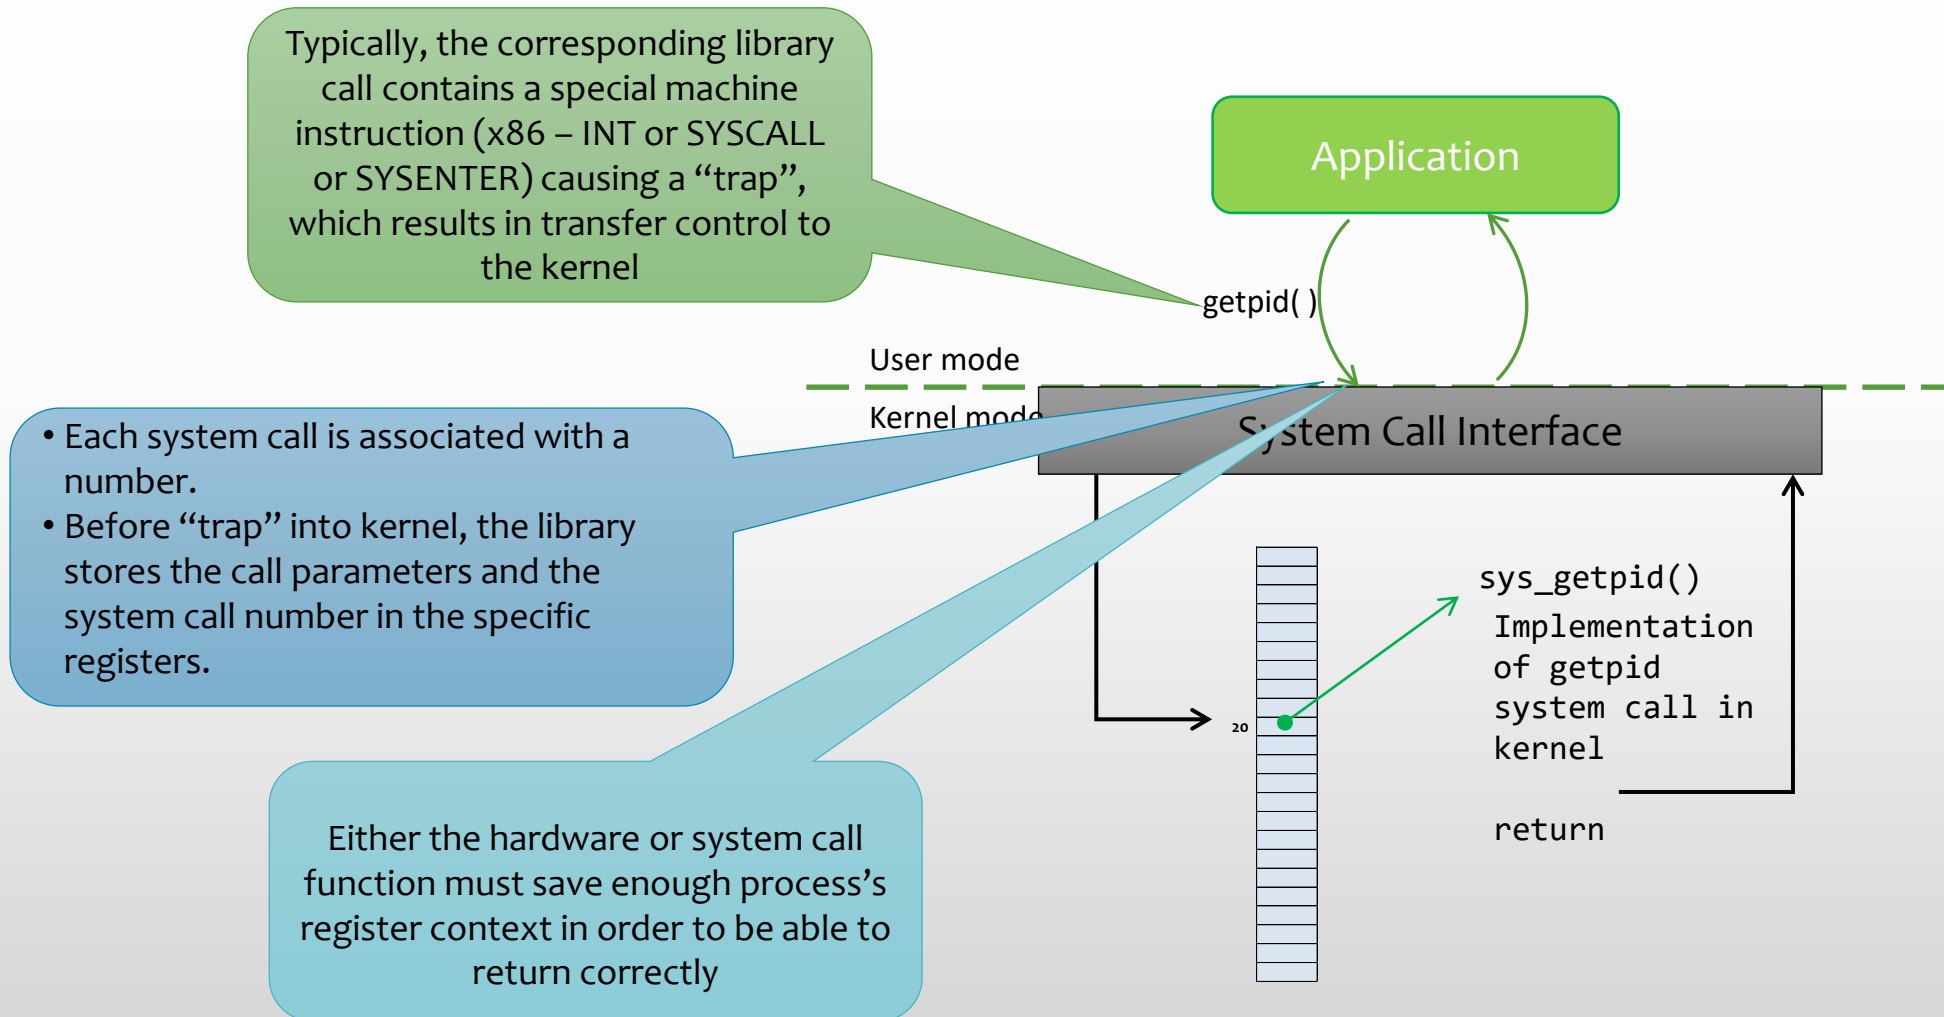
- System calls allow the kernel to carefully expose certain key services to applications
 - Most OSs expose a few hundreds such functions
- Applications mostly accessed via a high-level **Application Program Interface (API)** rather than directly invoke the specific system call
 - Common APIs are Windows API for Windows, and POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X)
 - Usually, in Unix system, the API is included in the run-time support library (e.g. C library)



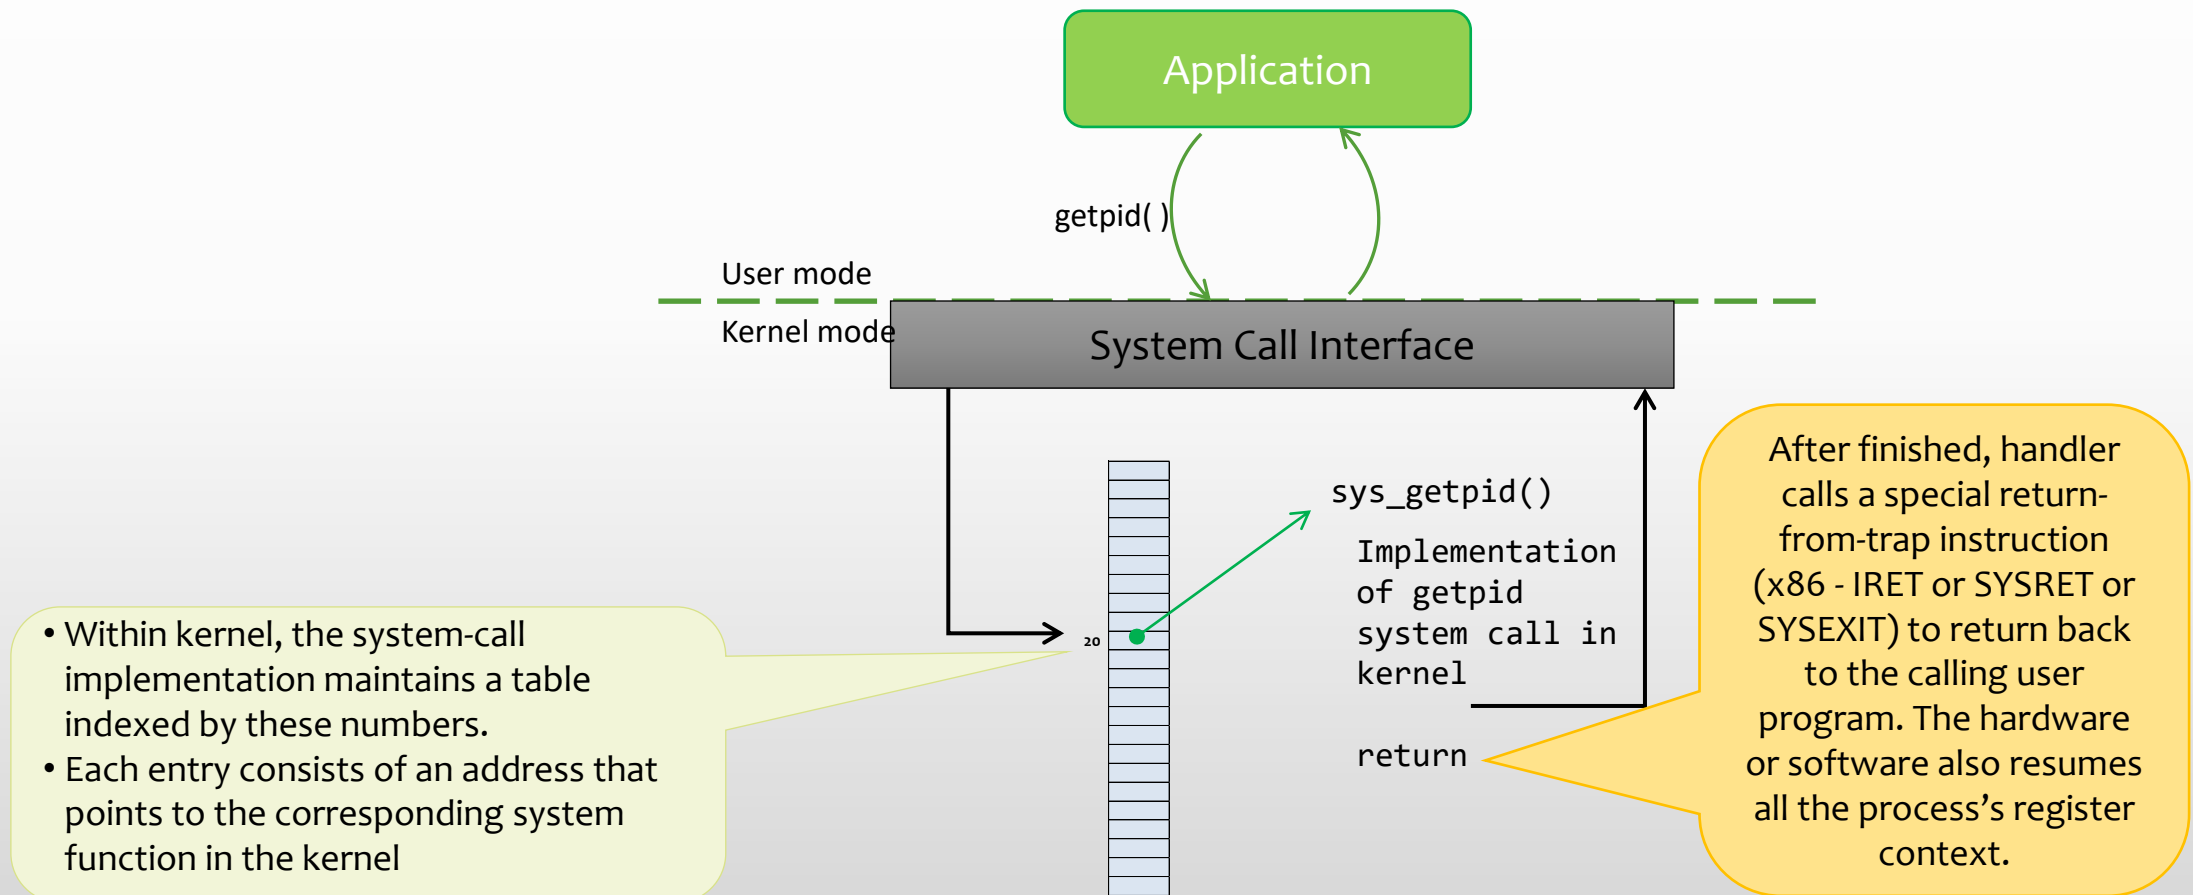
To request OS services

- Why use APIs rather than directly invoke system calls?
 - Caller does not need to know the details on how to invoke a system call under that OS & hardware
 - Different OSs use different techniques in implementing system calls
 - Usually, involves assembly code
- Programmer Just needs to know how to use the **standard** API to invoke the system call and understand what information will be returned by the system call via the API
 - Details of the OS interface are hidden from programmer by the API
 - The system call interface invokes intended system call in kernel and returns status and results of the system call and pass back to calling program via the API
 - The same API can be provided by different OS platforms → allows running the applications on different platforms

System Call Implementation



System Call Implementation



Virtualizing the CPU

- To provide illusion that each process has its own CPU
- By virtualizing, does not mean giving a virtual CPU
 - The program is directly running on the real CPU
 - Mechanism:
 - By running one process, then **stopping it and running another**, and so forth
- The Crux
 - How to **transparently** and **temporarily stop** a process and **resume** it?
 - With direct execution of application process on CPU, how can OS **regain control** of the CPU?

Switching from one process to another process

- To stop a process, one just needs to remove it from running to other state
- How can OS make sure that the process can be resumed **without affecting** its execution logic?
 - OS needs to **save the context** (e.g., CPU registers) of current running process **before stopping it**
 - To resume, OS needs to **restore the context** of the soon-to-be running process back to the CPU **before resuming it**
 - Question: Where to store the execution context of a process?
- **Context Switch**
 - Switching of the execution context of one process to another

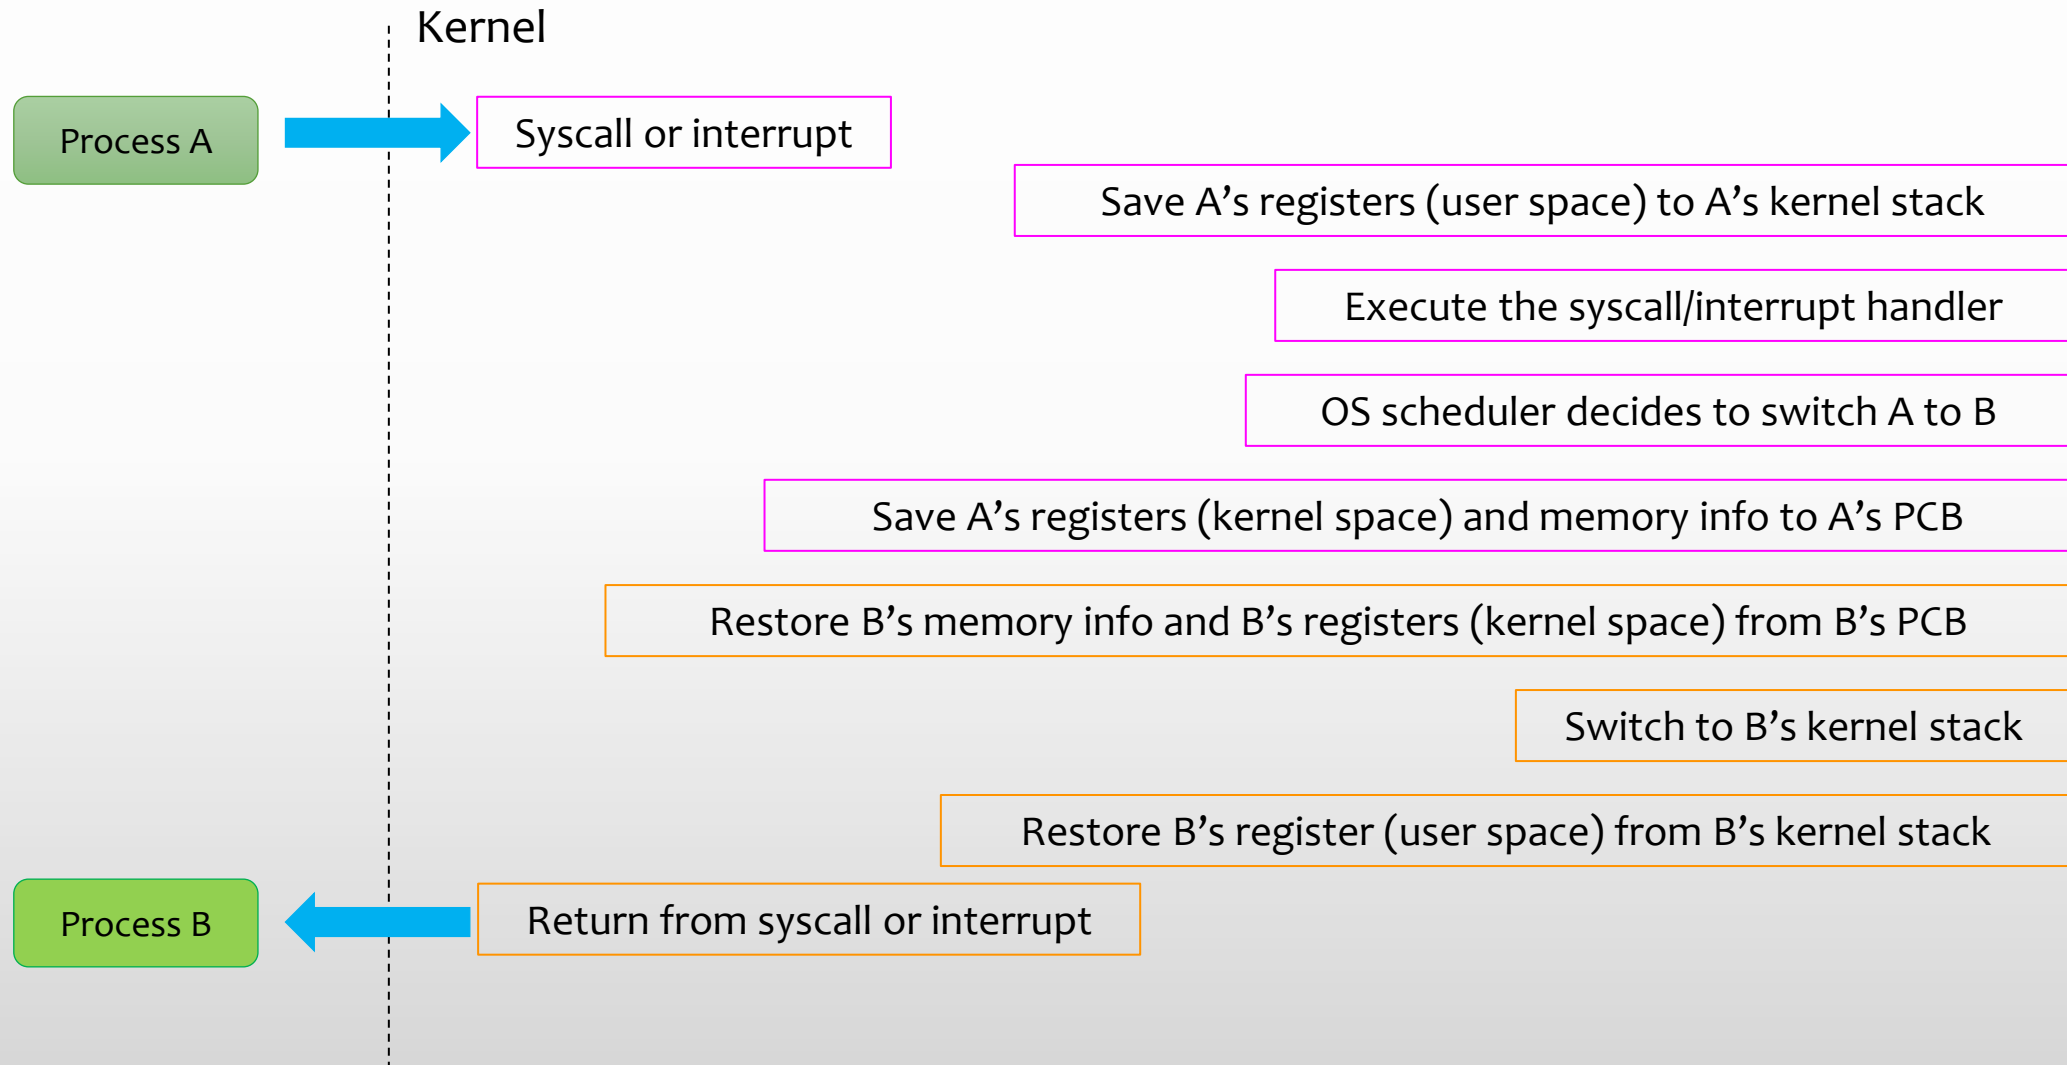
Mode Switch and Context Switch

- Mode switch
 - Switching from user mode to kernel mode
 - the Kernel is said to be “executing on behalf of the process”
 - the process’s context remains accessible (e.g., address space)
 - upon exiting kernel, the process *may* resume and return to execute in user space
- Context switch
 - Switching from one process to another
 - is an essential feature of multiprogramming or time-sharing systems
 - Kernel must
 - (1) suspends the progression of current process and stores its context
 - (2) retrieves the context of another process (B) and restores it to the CPU
 - (3) resumes execution of process B

Context Switch

- Action
 - Save the “current” process’s execution context
 - Change the process’s state in PCB to appropriate status
 - Move the process to appropriate queue
 - Select a ready process (according to **scheduling policy**)
 - Load the “to be dispatched” ready process’s execution context
 - Update the ”to be dispatched” process’s state of that going to be run process
 - Resume that process by turning control over to that process

Switching contexts



Context Switch Overhead

- The CPU is considered as not performing any “useful/real” computation for the application process
 - Thus, with frequent context switching, application would take longer time to complete its execution
 - OS must minimize context-switching time
 - There are hardware instructions in some architectures that facilitates the switching
- Has some indirect cost (induced overheads)
 - The newly scheduled process may not have data and/or instructions in the physical memory
 - The newly scheduled process does not have any part of its address space in the CPU cache

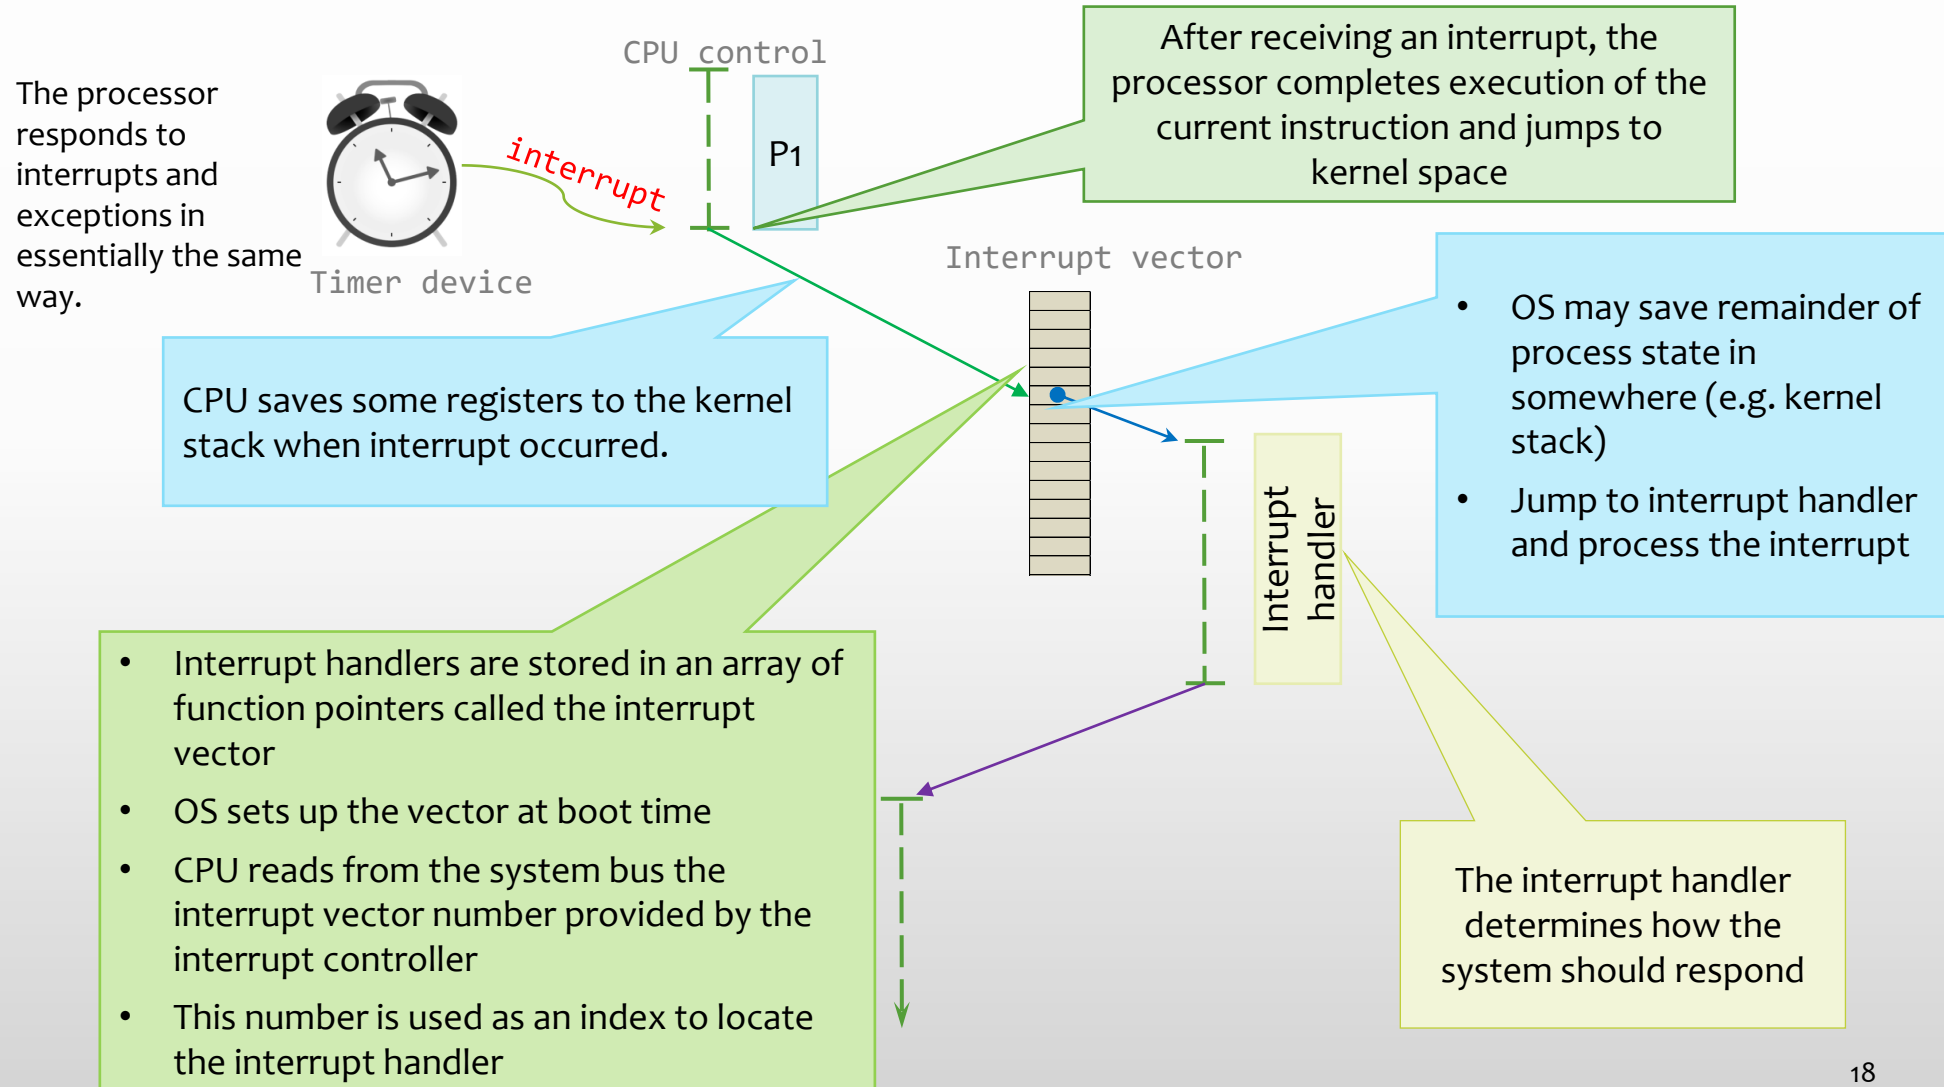
To regain control of CPU

- A Passive Approach
 - OS just waits for a process to make system calls or an illegal operation
 - Indeed, this works most of the time as processes quite frequently invoke system calls
 - However, in worst case, a process gets into an infinite loop and never makes a system call, then what can OS do!!
- An Active Approach (assume non-cooperative processes)
 - Needs hardware support
 - A timer device periodically generates an interrupt, e.g. every few milliseconds
 - When the interrupt is raised, OS will be called in to handle the interrupt, and thus can do what it wants

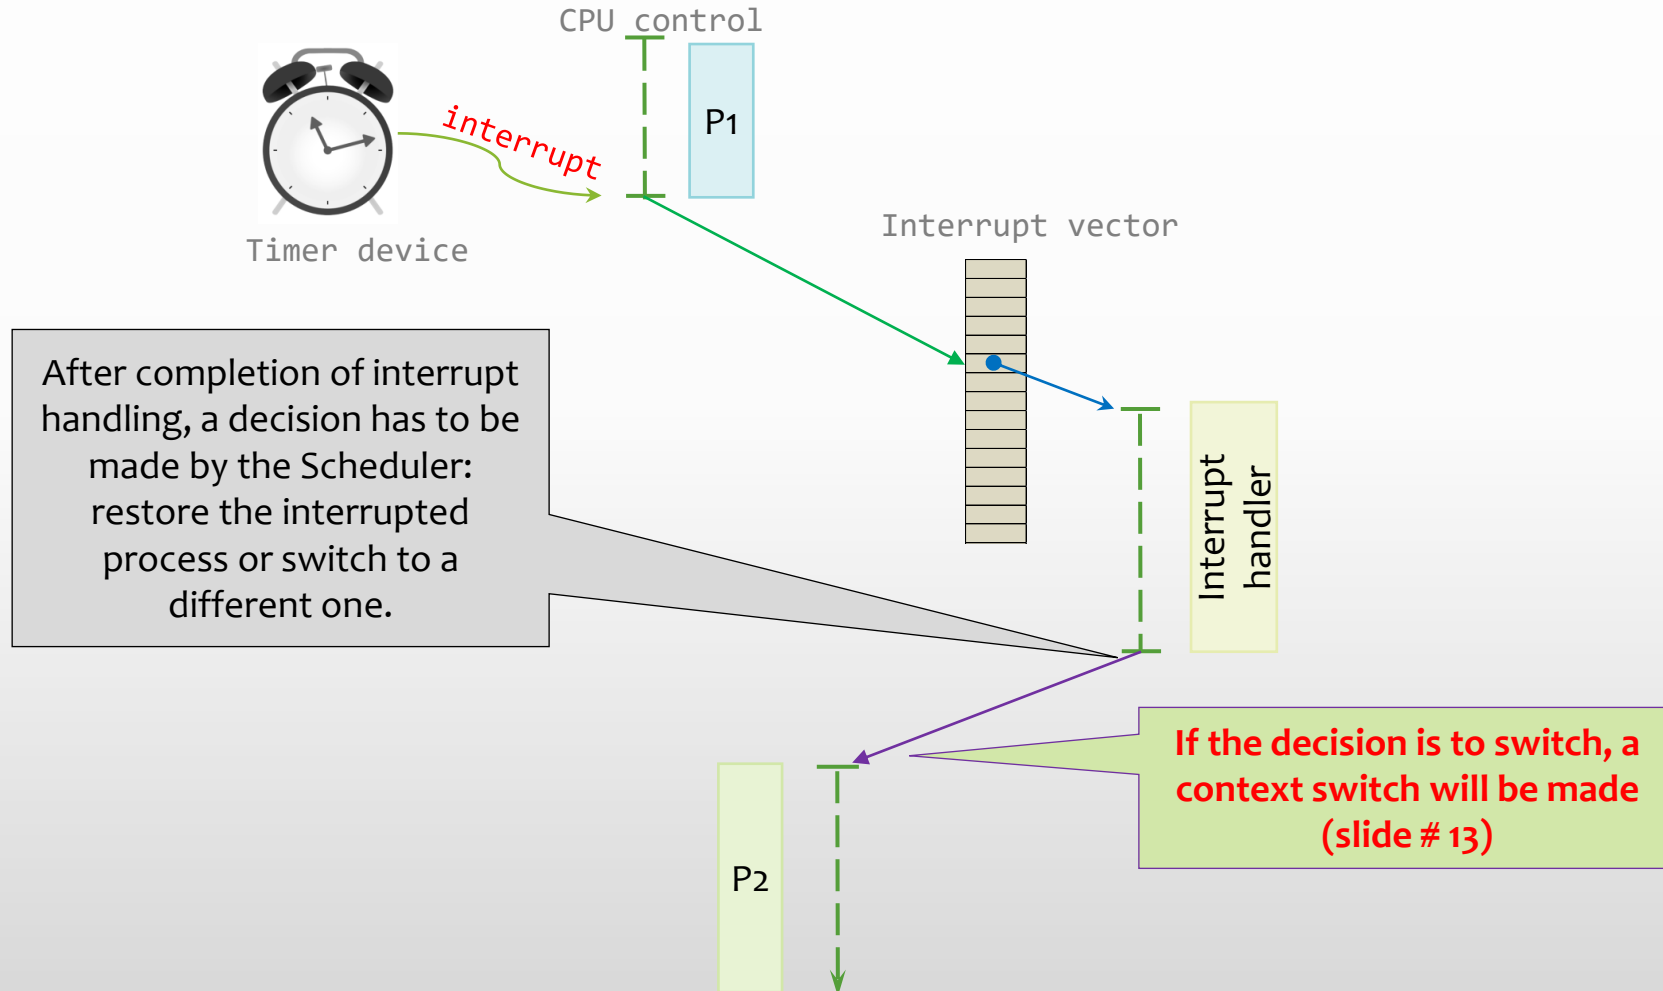
Interrupts

- An interrupt is an event raised by **software** or **hardware** when it needs the CPU's attention
 - Interrupts enable OS to respond to “**alerts**” immediately
- Hardware device applies signal to the interrupt line directly or via APIC to the CPU
 - e.g., a key is pressed or the mouse is moved
 - **Asynchronous** with the operation of the **current** running process
- May be **triggered by the current running process**
 - **Synchronous** with the operation of the **current** process
 - e.g., dividing by zero or referencing protected memory
 - Interrupt in this case is called an **exception**
 - Can be classified into 3 categories: Faults, Traps, & Aborts
 - Faults - CPU detects an issue before executing an instruction, e.g., page fault, segmentation fault, etc.
 - Some faults can be corrected and the interrupted program can resume
 - Traps – CPU detect the exception immediately following the execution of the instruction, e.g., INT instruction
 - After handling the trap, the interrupted program is allowed to continue
 - Abort – CPU experiences an unrecoverable error, e.g., double fault – another fault happens in the middle of handling previous fault.

Interrupt Handling



Interrupt Handling



Summary

- Three mechanisms
 - System call
 - Each call represent a service provided by the OS to user application
 - System mode will switch from user mode to kernel mode, and after the service, will switch back from kernel mode to user mode
 - Context switch
 - The key action of virtualizing the CPU
 - The whole set of register context must be saved for the interrupted process
 - The register context of coming process are restored
 - Context switch involves certain amount of performance overhead
 - Interrupt processing
 - A mechanism for hardware devices to alert OS for handling high-priority events
 - This gives a chance to OS to regain control of CPUs