

Thread Abstraction & Concurrency

2023/24 COMP3230B

Concurrency

- “Why are we studying this in OS class?”
 - “History” is the one-word answer.
 - Simply put, the **OS was the first concurrent program**, and thus most of these techniques arose due to the need for them within the OS. Later, as multi-threaded programs became popular, application programmers also had to consider such things.

Contents

- What are threads?
 - Why multithreading?
- POSIX Threads (Pthreads)
- Concurrency Issue
 - Race Condition
 - Critical Section
 - Mutual Exclusion
 - Synchronization

Related Learning Outcomes

- ILO 2a - **explain** how OS manages processes/**threads**
- ILO 2c - **explain** the underlying causes of **concurrency** issues
- ILO 4 - demonstrate knowledge in applying system software and tools available in modern operating system for software development

Readings & References

- Required Reading
 - Chapter 26 – **Concurrency: An Introduction**
 - <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-intro.pdf>
 - Chapter 27 – **Interlude: Thread API**
 - <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-api.pdf>

What are threads?

- Thread of execution
 - A **sequence of instructions** that performs a task (within the application process)
 - We can view traditional process as a process with one thread of execution
- Multithreaded process
 - A process has multiple threads of execution
 - A thread is an entity **within** a process
 - Multiple threads within a process
 - **can execute concurrently**
 - **share the process's address space** and other global info

What are threads?

```
main() {  
    funcA()  
    funcB()  
    funcC()  
    funcD()  
}
```

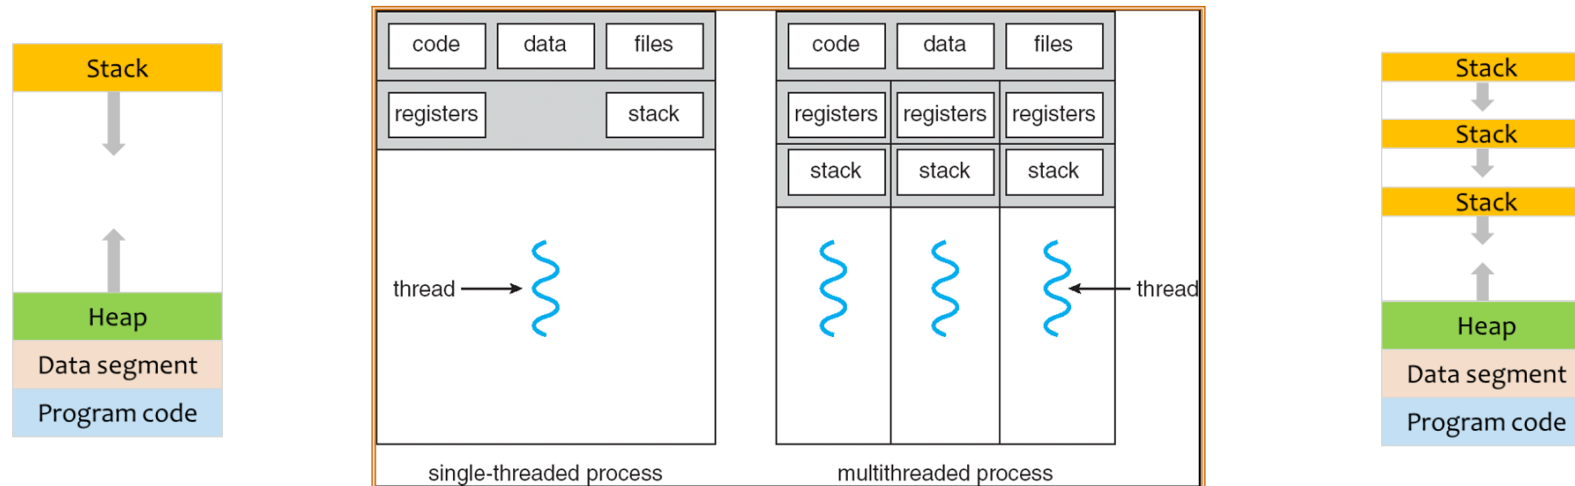
```
main() {  
    funcA()  
    funcC()  
    write()  
}
```

```
main() {  
    funcB()  
    funcD()  
    read()  
}
```

```
Thread1 {  
    funcA()  
    funcC()  
}  
Thread2 {  
    funcB()  
    funcD()  
}  
  
main() {  
    Thread1  
    Thread2  
    Join()  
}
```

What are threads?

- To support multiple threads within a process, the system must provide each thread with its own
 - program counter; private set of registers
 - stack (& kernel stack)
 - and its own control block – **Thread Control Block (TCB)**
- Like process, each thread transits among a series of discrete thread states: new, running, ready, blocked, and terminated



Why Multithreading?

- Make inherently **parallel tasks** simpler to express in code
- For performance
 - Less set up is needed, it takes **less time to create/terminate** a thread than a process
 - A process with multiple threads could continue running even if one of its threads is blocked
 - A traditional process will be moved to blocked queue if it calls a blocking system call
 - Can make use of underlying multicores
 - One process with multiple threads; each thread can run on a core
 - **Less overhead in switching** between threads of the same process
 - thread switching as compared to process switching

POSIX Threads - Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- Why use POSIX threads?
 - Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- API specifies behavior of the thread library, implementation is up to development of the library
 - POSIX states that processor registers, stack and signal mask are maintained individually for each thread
 - POSIX specifies how operating systems should deliver signals to pthreads in addition to specifying several thread-cancellation modes

Thread Creation

- `pthread_create()` creates a new thread and makes it executable.
 - Typically, threads are first created from within `main()` inside a single process.

```
#include <pthread.h>
int pthread_create(pthread_t * thread,
                  const pthread_attr_t attr,
                  void *(*thrfunc)(void *),
                  void *args);
```

- 4 arguments:
 - **thread**: a pointer to a structure of type `pthread_t`, which becomes the handler of a thread
 - we can pass this handler to various thread operations
 - **attr**: used to set thread attributes
 - in most case, **set it to NULL** to use the default setting
 - **thrfunc**: a function pointer, which **points to the C function** that the **thread will start executing** once it is created
 - **args**: a **single** argument to be passed to the **thrfunc** function.

Thread Termination

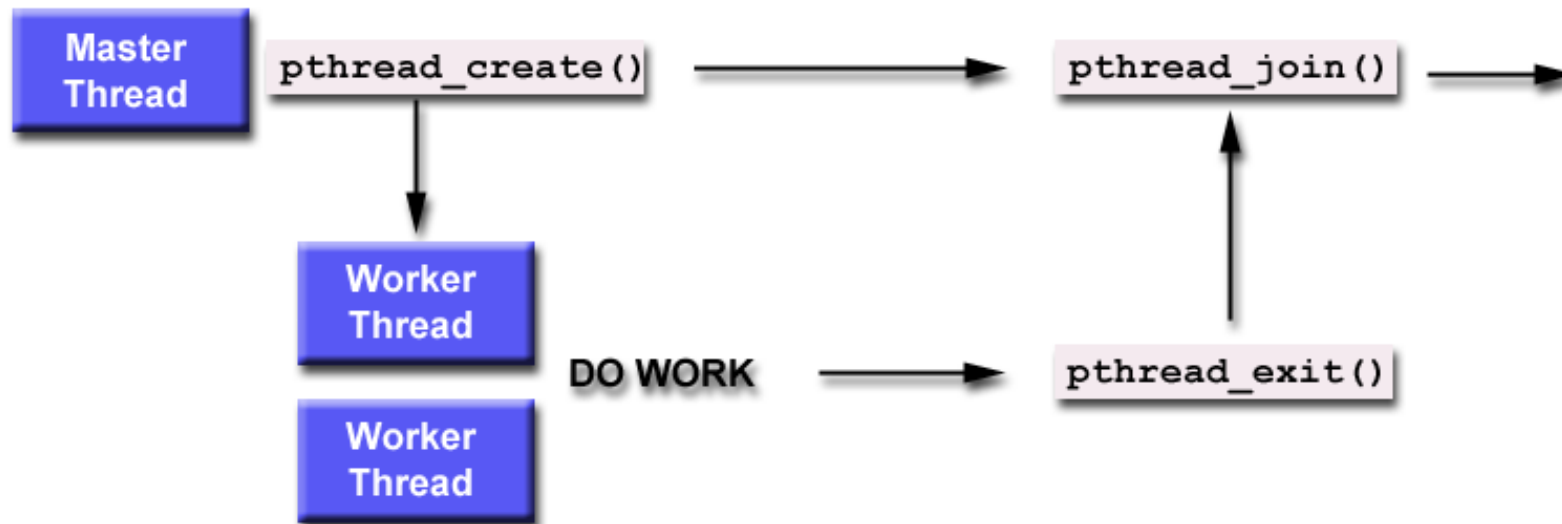
- `pthread_exit()` terminates the execution of the calling thread
 - Typically, a thread calls this after it has completed its work

```
#include <pthread.h>
void pthread_exit (void *retval)
```

- `retval`: return value of the thread
 - It can be retrieved by another thread using `pthread_join()`
- It **does not close files**; any files opened inside the thread will remain open after the thread is terminated

Waiting for a thread to terminate

- “Joining” is one way for a thread (especially main/master thread) to **wait for** other threads to exit. For example:



Thread Join

- The pthread_join() **blocks the calling thread until** the specific thread terminates

```
#include <pthread.h>
int pthread_join (pthread_t thread_id,
                  void ** value_ptr);
```

- **thread_id**: specify **which thread** in which this calling thread will **wait for**
- **value_ptr**: a return value you expect to get back from the target thread (via a pointer to void pointer)
 - If you don't care, set it to NULL

An Example

```
#include <stdio.h>
#include <pthread.h>

void *func1 (void *arg){
    int x = *((int*)arg);
    printf("The integer passed in is %d\n", x);
    printf("Thread: Process id is %d\n", (int)getpid());
    pthread_exit(NULL);
}

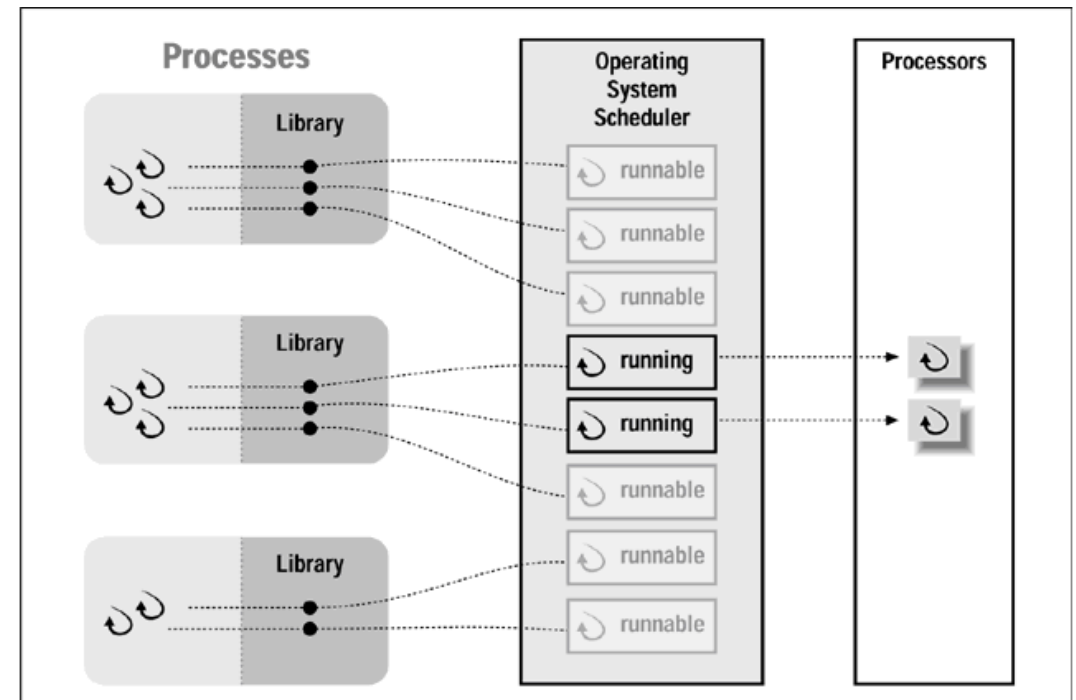
int main() {
    pthread_t thread_id;
    int x = 1;

    printf("Main process: Process id is %d\n", (int)getpid());

    pthread_create(&thread_id, NULL, func1, (void*)&x);
    pthread_join(thread_id, NULL);
    return 0;
}
```

Pthreads in Linux


- In Linux, each pthread is implemented as standard process – lightweight process (LWP)
 - All LWPs in the same multithreaded application share the memory address space, the open files, global variables, heap, ...
 - To enable threading, Linux uses the clone() instead of fork()
 - Clone accepts arguments that specify which resources to share with the child task
- It means the scheduler does not differentiate between a thread and a process



Concurrency

Concurrency Issue


- Suppose the value of x was originally 10, what will be the outcome after two threads executed the instruction?



- It can be 11 or 12
- Why? And how can we guarantee to always get 12?

Reasons

- Multiprocessor or multi cores
 - The 3 instructions can be executed by 2 cores at the same time
- Single core



TS – thread switch

Concurrency Issue

- The main issue is **uncontrolled scheduling**
 - When will a thread be executing is not predictable
 - We cannot predict **at what time** and **for how long** a thread is being schedule to run
- **Race Condition**
 - We called the scenario that **several threads access and modify** a **shared data item concurrently** and the outcome of the execution depends on the particular order (**a race**) in which the accesses took place
 - This results in non-deterministic computation, where it is **not known what the result** will be and it is *indeed likely to be different across runs*

Critical Sections

- For multithreaded program, most code is safe to run concurrently
 - When not accessing and modifying shared data
- Blocks of code where a particular shared data is modified must be guarded
- We called these blocks of code – **Critical Sections**
 - We would like to have only one thread be in its critical section accessing the specific protected shared data at one time
 - it should execute as quickly as possible

Mutual Exclusion

- Therefore, shared data in critical sections must be accessed in mutually exclusive way
 - Only one thread is allowed access at one time
 - The winner thread is in-effect “locked” the shared data
 - Others must wait until the shared data is unlocked
- This is called serialized access or Mutual Exclusion

Atomic Operation

- One way to solve the race condition in “x++” is to make the operation to be **ATOMIC** –
 - **The operation cannot be interrupted in the middle**, and hardware can guarantee either has done or not done at all
- Unfortunately, critical sections may contain more complicated operations, e.g., linked list traversal
- Just having Atomic operations is not good enough

Synchronization

- Within a multithreaded process, there is another common interaction
 - One thread has to wait for another thread to work on some action before it continues
 - e.g., a worker has to wait for a task assigned by the boss
- The Crux
 - How to provide support for mutual exclusion? for synchronization? What support do we need from hardware and the OS?

Summary

- Threads have another name – light-weighted processes.
 - Each thread is a code fragment, within a process, that can be scheduled and executed independently
 - A thread has its own program counter, registers' contents, and stack, but it shares the same process's address space with other threads
 - Similar to process management, OS uses a thread control block to abstract a thread entity
- With multiple processes or threads, we have to face the concurrency issues
 - One of the responsibilities of OS is to provide mechanisms for processes/threads to synchronize and coordinate between processes/threads