

# Synchronization Tools

2023/24 COMP3230B

# Contents

- mutual exclusion – mutex lock
- Using atomic instructions to build spin-waiting mutex lock
- Condition variables – a synchronization tool
- Semaphores – another tool that can serve for locking and/or synchronization

## Related Learning Outcomes

- ILO 2c – describe the principles and **techniques** used by OS to support concurrency and synchronization control.
- ILO 4 – [**Practicability**] demonstrate knowledge in applying system software and tools available in modern operating system

# Readings & References

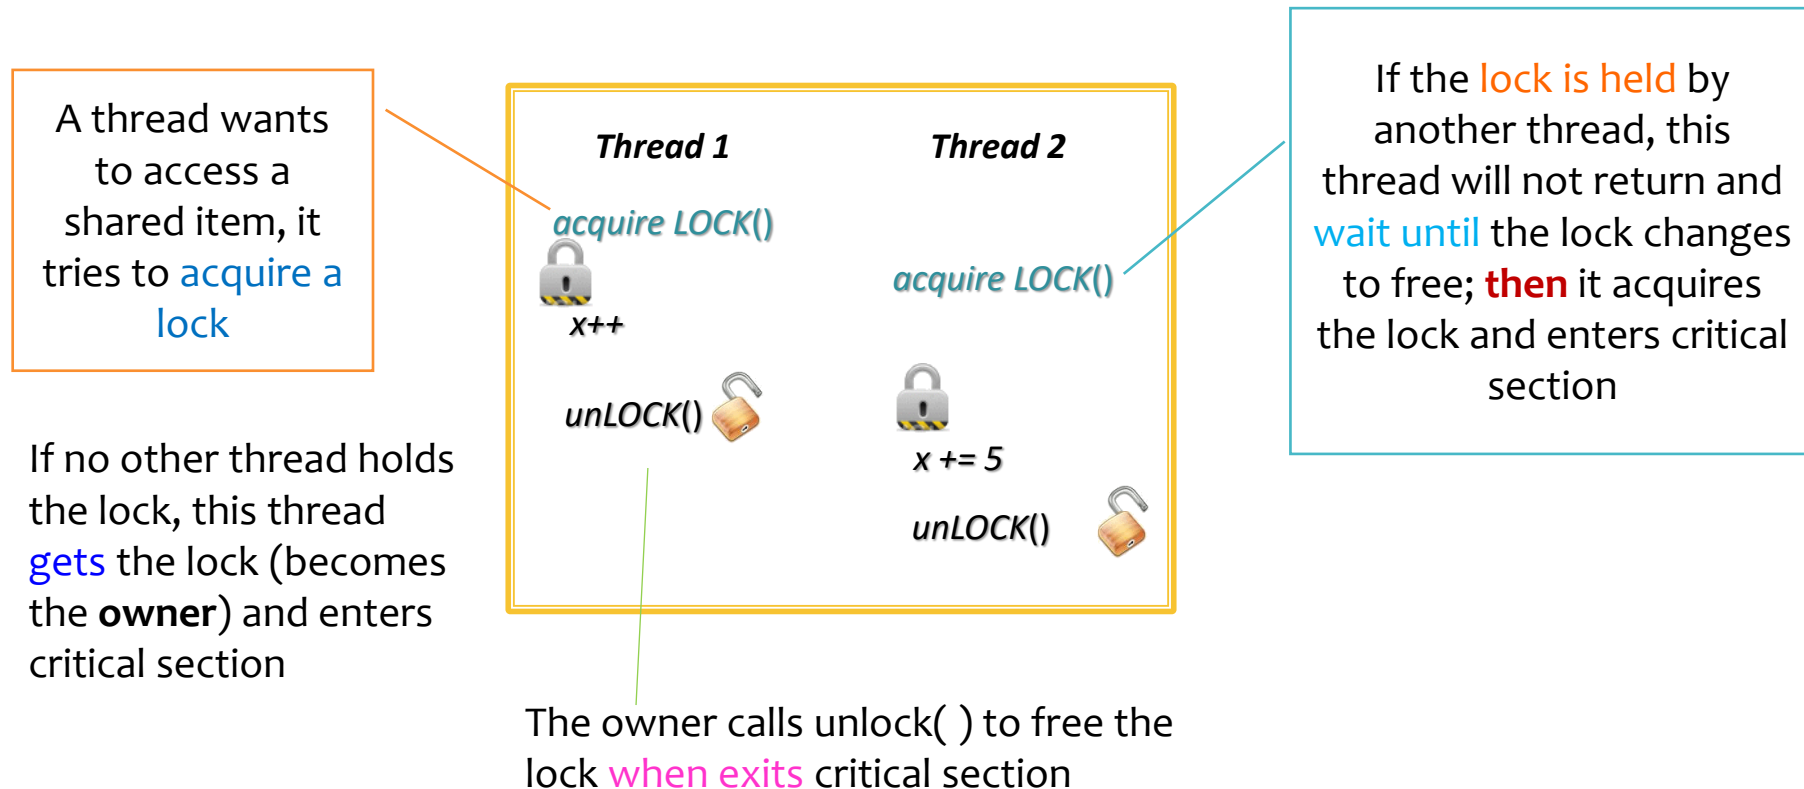
- Required Readings
  - Chapter 28 – **Locks** (except sections 28.10, 28.11, 28.15, 28.16)
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-locks.pdf>
  - Chapter 30 – **Condition Variables**
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-cv.pdf>
  - Chapter 31 – **Semaphores**
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/threads-sema.pdf>

# Mutual Exclusion

- We need a mechanism to inform OS that a thread is **going to enter** its critical section
- Here comes the concept of using a “lock”
  - A lock is a data structure used for **indicating** the start and end of **a critical section**
    - Indicate shared data is about to be accessed and **ask for the system** to provide necessary protection

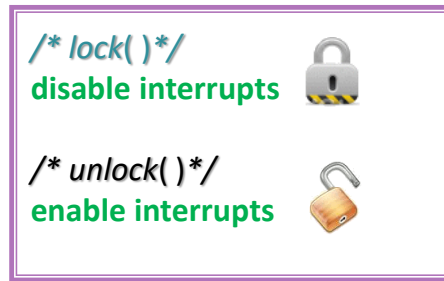
# Mutual Exclusion – Lock

- A lock is either available (free or not locked) or held (locked or acquired)



# Building Locks

- The main cause of the indeterministic outcome is **uncontrolled scheduling**
  - Can we avoid being preempted when the thread is in critical section?



- **Disabling interrupts**
  - Prevents current executing thread from being preempted, as **without interrupts, scheduler will not be invoked**; thus, no other threads will be able to turn to running state
- Works **only** on systems with **single core**
- Being used rarely, mostly used in kernel **under privilege mode**

# Atomic Instructions

- Modern machines provide special atomic hardware instructions
- **Test-and-Set** Instruction
  - It enables us to test the old value **while simultaneously set** the variable to a new value
    - The instruction **returns the old value** and **simultaneously updates the variable** to **new value** in an **atomic way**
- **Compare-and-Swap** Instruction
  - Test whether the content of a variable is equal to expected; if so, update the variable to a new value; otherwise, do nothing
    - The instruction **always returns the current value** of the variable

Guarantee execution till end without interruption

```
int TestAndSet(int *ptr, int new) {  
    int old = *ptr;           //get current value  
    *ptr = new;               //set to new value  
    return old;               //return old value  
}
```

Guarantee execution till end without interruption

```
void CompareAndSwap(int *ptr, int expected, int new)  
{  
    int temp = *ptr;          //get current value  
    if (temp == expected)  
        *ptr = new;  
    return temp;  
}
```



# Test-and-Set

- Implementing Mutual Exclusion
  - Based on **checking** of the **value of a shared variable** to decide whether a process can enter a Critical Section
    - Uses **busy waiting** (spin-wait or spin-lock) to test whether it can enter the Critical Section

If another thread is in critical section, flag is 1; in that case, TestAndSet( ) will return 1

If no thread is in critical section, flag is 0; in that case, TestAndSet( ) will return 0 and will atomically modify flag to 1

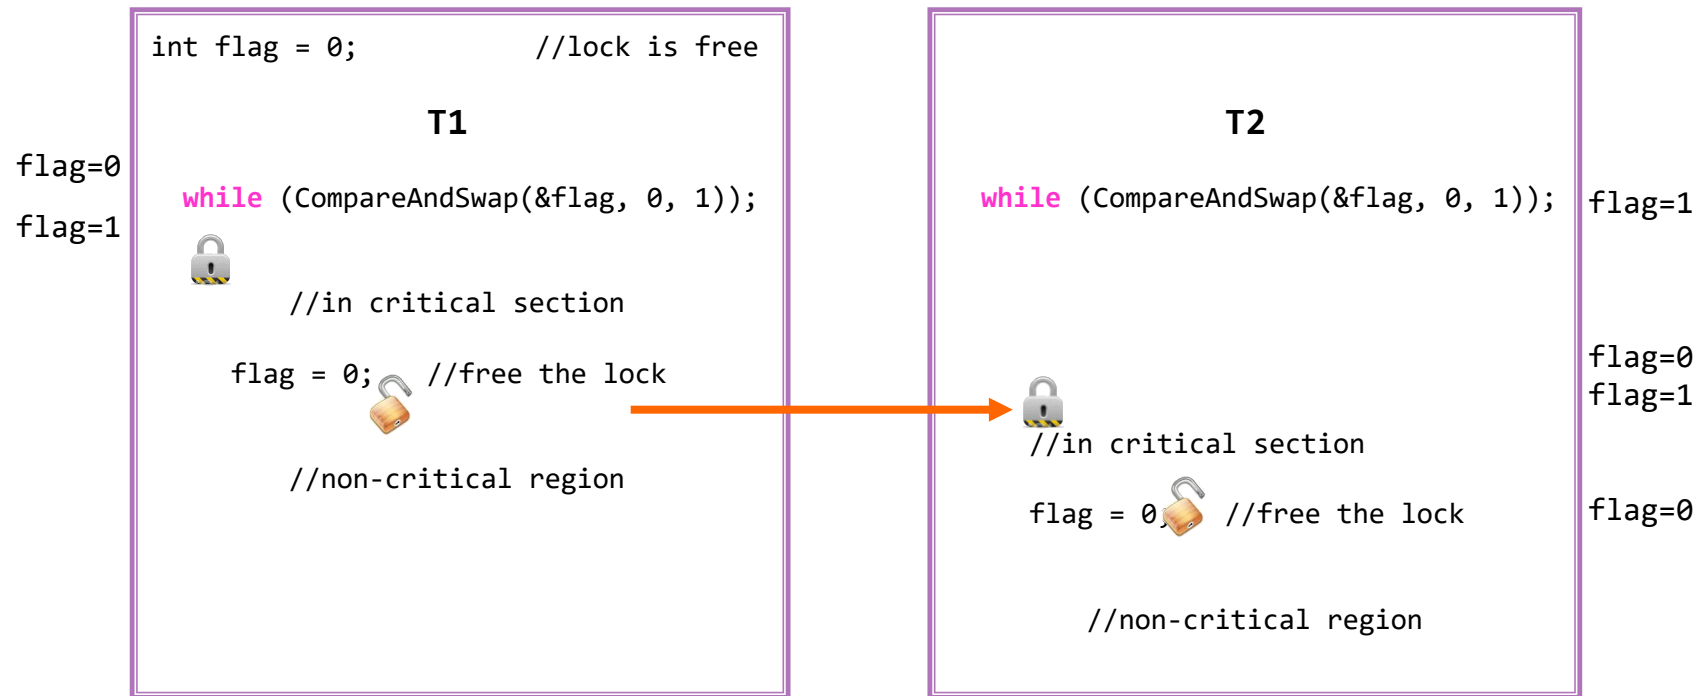
```
int flag = 0;           //when flag = 0, lock is free
while (TestAndSet(&flag, 1) == 1);    //lock( )
// into the critical section
flag = 0;               //unlock( )
```

# Test-and-Set

- With the guarantee of atomic update of a shared variable
  - This guarantees mutual exclusion as **only one thread can set the shared flag to 1**, others have to spin-wait for the flag to turn back to 0
- Spin-wait / Busy-waiting
  - However, the processor is essentially doing nothing but just executes the while() loop
  - **Wastes significant processor time**
    - those threads may just spin-waiting until the time quantum expires
- Spin-waiting cannot avoid indefinite postponement; thus, **do not provide fairness guarantee**
  - e.g., when two threads T1 & T2 contend for entering to critical sections, because of uncontrolled scheduling, it is possible that T1 always be the lucky one & T2 always finds that the flag is always 1

# CompareAndSwap

- Example: Intel IA-32 and IA-64 contain an CMPXCHG instruction
- Implementing Mutual Exclusion
  - Again, using spin-waiting



# Mutex – Hardware Instructions

- Atomic instructions
  - Applicable to any number of processes/threads on either a single processor or multiple processors sharing main memory
- Using Machine instructions cannot provide a good mutual exclusion solution on its alone
  - For example, possibility of having indefinite postponement if more than one thread is waiting
- Unfortunately, Spin-waiting consumes processor time
- The Crux
  - How can we develop a lock that **does not needlessly waste time spinning** on the CPU as well as **maintain fairness**?
- The answer is: we need OS support.

# Pthread Lock


- In POSIX library, the lock data type is called mutex
- mutex variables must be declared with type
  - `pthread_mutex_t`
- Must be initialized before they can be used
- Two operations on mutex variables
  - `pthread_mutex_lock()` and `pthread_mutex_unlock()`
    - **OS blocks the calling thread** if the request lock is not available
      - So the thread transits to Blocked STATE
- When finished using a mutex, deletes it with
  - `pthread_mutex_destroy(&lock);`

# Control Synchronization

- Another typical interaction between threads is the control synchronization
  - a thread (A) should **perform** some action **only after** some other threads **have performed** specific actions (or **have detected** an event/condition)
  - **Thread A has to be waited and some other thread has to notify A** about the occurrence of the event/condition
- A simple approach – Thread A **just spins until the condition** becomes true
  - this is inefficient and wastes CPU cycles.

# Condition Variables

- A data type explicitly designed to support synchronization between threads (without wasting CPU cycles)
  - it has a **queue** that threads can wait over there when some condition/event is not met
  - **some other thread**, once detects the condition/event, **inform (signal)** those waiting threads **and wake one** of those waiting threads
  - when using condition variables, each condition variable can be associated to a distinct condition/event






# Pthread Condition Variables

- Must be declared with type `pthread_cond_t`, and must be initialized before they can be used
- To wait for a condition
  - `pthread_cond_wait()`
    - When call this function, the calling thread is BLOCKED until wake up by other using `pthread_cond_signal()/broadcast()`
- To inform others that the condition/event has happened
  - `pthread_cond_signal()`
    - Wake up one thread which is waiting on the condition variable
  - `pthread_cond_broadcast()`
    - Wake up all threads currently blocked on the condition variable
- Finish using a condition variable
  - `pthread_cond_destroy(&cond);`

# Producer-consumer problem

- Classical synchronization problem, also known as Bounded-Buffer Problem
  - The system has a **finite buffer pool** shared by one or more producers and one or more consumers
    - Each buffer is capable of holding one unit of information
  - Producers produce data items and wish to place them in buffers
    - A producer thread produces a data item and places it into the next available buffer in the buffer pool
  - Consumers take out data items from buffer and consume/process the data item in some way
    - A consumer thread consumes a data item by removing it out of the buffer pool



# Producer-consumer problem

- Synchronization requirements
  - Buffer pool is shared resource, producers and consumers need to use some method to **coordinate the access** to the pool
  - A producer must **not overwrite** a buffer when buffer pool is full
  - A consumer must **not consume** an empty buffer when buffer pool is empty
  - Mutual exclusion
  - Information must be consumed in FIFO order

# A Solution to Producer-consumer problem

To wait for  
buffer pool  
becomes  
NOT FULL

```
buffer[sizeofbuffer];  
mutex_t mv;  
cond_t notFULL, notEMPTY;  
count = 0;
```

To wait for  
data/job in buffer  
pool

```
Producer( ){  
    while (1)  
    {
```

Buffer is  
definitely  
NOT FULL

```
        d = generating_data();  
        lock(mv);  
        while (count == sizeofbuffer)  
            cond_wait(notFULL, mv);  
        append(d, buffer);  
        cond_signal(notEMPTY);  
        unlock(mv);  
        remaining_work();
```

```
//try acquire the mutex lock  
//while buffer pool is full  
//wait for free space  
//update the count  
//data is available in buffer  
//release the lock
```

Buffer pool is definitely  
NOT EMPTY.  
Inform waiting threads if  
any

- Before waiting, it releases the mv lock.
- Then queues in the cond queue.
- Upon waking up, it gets back the mv lock before return

# A Solution to Producer-consumer problem

```
Consumer( ) {  
    while (1)  
    {
```

Buffer is  
definitely NOT  
EMPTY

```
        lock(mv);  
        while (count == 0)  
            cond_wait(notEMPTY, mv);  
        d = take(buffer);  
        cond_signal(notFULL);  
        unlock(mv);  
        consume_data(d);  
        remaining_work();  
    }
```

//try acquire the mutex lock  
//while buffer pool is empty  
//wait for data item  
//update the count  
//one more free buffer space  
//release the lock

Buffer pool is definitely  
NOT FULL. }  
Inform waiting threads  
if any

- Before waiting, it releases the mv lock.
- Then queues in the cond queue.
- Upon waking up, it gets back the mv lock before return

# Semaphores

- Another synchronization tool that can serve as **mutex locks** or use for **control synchronization**
- A semaphore is an object with an **internal protected integer** variable
- Must be declared with type **sem\_t**
- The **integer value** stored in the semaphore **determines its behavior**
- To initialize semaphore, **sem\_t k**, to a value of **x**
  - `sem_init(&k, 0, x);`

## Semaphores (2)

- The protect integer value **can only be** accessed via two **atomic** operations
  - `sem_wait(&k)` and `sem_post(&k)`
  - Also called **P()** (or **wait** or down) and **V()** (or **signal** or up), respectively

### `sem_wait(&k)`

```
sem_wait(&k) {  
    decrement the value of semaphore k by 1  
  
    if value of semaphore k < 0  
        place the thread in k's waiting queue  
}
```

### `sem_post(&k)`

```
sem_post(&k) {  
    increment the value of semaphore k by 1  
  
    if any threads are waiting on k's queue  
        wake up one thread  
}
```

# MutEx – Binary Semaphores

- Semaphores that work like a mutex lock
- Binary – semaphore can have the value **one or zero only**
  - allow only one thread in its critical section at once
  - sem\_t lock is **initially set to one only**

```
sem_t lock;  
sem_init(&lock, 0, 1);
```

Thread 1

```
while (1) {  
    sem_wait(&lock);  
    //in critical section  
    sem_post(&lock);  
    //other code fragment  
}
```

**lock = 1**  
**lock = 0**

Thread 2

```
while (1) {  
    sem_wait(&lock);  
    //in critical section  
    sem_post(&lock);  
    //other code fragment  
}
```

**lock = 0 (q=1)**  
**lock = 0 (q=1)**

Thread 3

```
while (1) {  
    sem_wait(&lock);  
    //in critical section  
    sem_post(&lock);  
}
```

**lock = 0 (q=2)**  
**lock = 0**  
**lock = 1**



# Readers-Writers Problem

- Another classic synchronization problem
- Consists of a set of threads accessing some shared data
  - Readers – threads that only reads the data
  - Writers – threads that modifies the data
- Typical example – Database access, linked list update and lookup
- Synchronization requirements
  - Many readers can perform reading concurrently
  - Reading is prohibited while a writer is updating
  - Only one writer can perform updating at any time

# A Solution using Binary Semaphore

```
Semaphore semMutex = 1;  
Semaphore readMutex = 1;  
int readcount = 0;
```

```
Reader() {  
    StartRead();  
    do_reading();  
    EndRead();  
}
```

```
Writer() {  
    StartWrite();  
    do_writing();  
    EndWrite();  
}
```

```
StartRead() {  
    P(readMutex); //Get the readMutex  
    readcount++; //Count the number of readers  
    if (readcount == 1) //If is the first reader  
        P(semMutex); //Get the access right  
    V(readMutex); //Release, others can update readcount  
}
```

```
EndRead() {  
    P(readMutex); //Finish reading, try decrement readcount  
    readcount--;  
    if (readcount == 0) //If is last reader, free access semaphore  
        V(semMutex);  
    V(readMutex);  
}
```

```
StartWrite() {  
    P(semMutex);  
}
```

```
EndWrite() {  
    V(semMutex);  
}
```

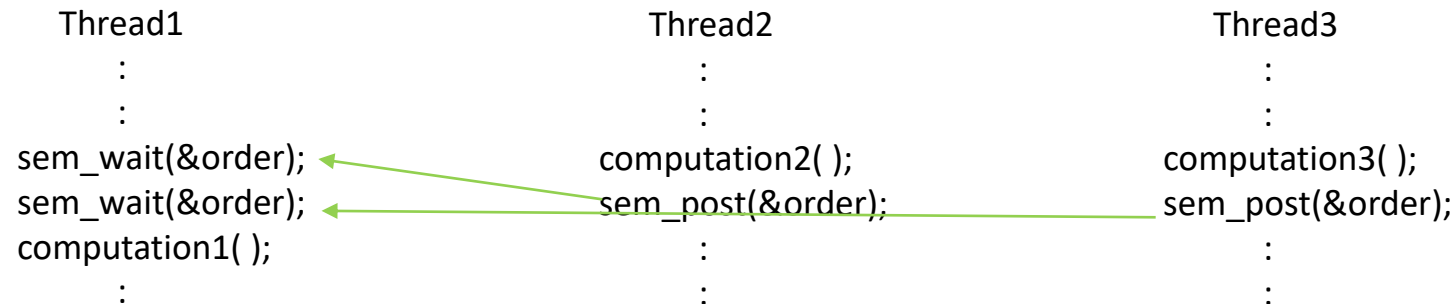
# General Semaphores

- Also known as **Counting Semaphores**
- Can be used to control access to a pool of **identical** resources
  - Initialized with **sem\_t** pool to the resource count
  - Decrement the semaphore when taking resource from pool – **sem\_wait(&pool)** operation
  - Increment the semaphore when returning it to pool – **sem\_post(&pool)** operation
  - If no resources are available, thread is blocked until a resource becomes available

# Synchronization bet. threads

- Semaphores can be used to notify other threads that specific event/condition have occurred/reached
- Example: Another example of common synchronization problems is **order of execution** of operations of different threads
  - Thread1 wants to wait for both Thread2 and Thread3 to finish their tasks before its operation

```
sem_t order;  
sem_init(&order, 0, 0); //initialize to 0
```



# Summary

- To avoid race condition, such that only one thread can be in critical section, we need to use mutual exclusion primitives, such as, mutex lock or binary semaphore, as a guard for accessing/modifying the shared data
- For spin-wait locks, in uniprocessor, because of the busy waiting, performance overhead can be quite painful
- By the use of condition variables, thread can have a non-busy waiting mechanism, by put itself into sleep and wait for other thread to wake it when the desire condition is met
- Semaphores are a powerful and flexible primitive that support both synchronization and mutual execution