



# Mechanisms: Segmentation & Paging

2023/24 COMP3230B

# Contents

- Segmentation
  - Address Translation
  - Protection and Sharing of Memory
  - External Fragmentation
  - Strategies in Managing Free Space
    - Best-fit, Worst-fit, First-fit, & Next-fit
- Paging
  - Virtual Page and Page Frame
  - Address Translation using Page Table

# Related Learning Outcomes

- ILO 2b - describe the principles and techniques used by OS in effectively virtualizing memory resources.
- ILO 3 [Performance] - analyze and evaluate the algorithms of . . . and explain the major performance issues . . .

# Readings & References

- Required Readings
  - Chapter 16 – Segmentation
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/vm-segmentation.pdf>
  - Section 17.1 & 17.3 of Chapter 17 Free-Space Management
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/vm-freespace.pdf>
  - Chapter 18 – Paging: Introduction
    - <http://pages.cs.wisc.edu/~remzi/OSTEP/vm-paging.pdf>

# Segmentations

- A segment is a **block of contiguous virtual addresses** of a particular length within a process's address space, which is **for a specific purpose/usage**
  - e.g., text segment, data segment, heap segment, and stack segment
- OS can place **each segment of a process** in **different parts** of physical memory
- Only allocate physical memory to memory segments that are in use; thus, **reduce** internal fragmentation → avoid waste
- **Assumption 2** (in Lecture 8 slide # 14) is no longer valid and **Assumption 3** is not relevant anymore

# Hardware Support

- To support segmentation, MMU needs to know how to translate **multiple logical segments** **within** a process's address space
  - **Multiple** base-and-bounds register pairs
  - **Each** base-and-bounds pair for the **address translation of one logical segment**

- Address Translation of a virtual address  $v'$

offset = virtual address  $v'$  – starting virtual address of **the segment**


physical address = segment\_base + offset

Starting physical address of the segment

- Protection

Upper limit of the segment

if offset > segment\_bounds then generate **segmentation fault**




# Operational Issue

- Given a virtual address  $v'$ , **which segment is this address in?**
  - We need to know **which segment** this virtual address is in so as to find the corresponding base-and-bounds pair
- Each virtual address  $v'$  is in one segment only. **Can we make use of the information within a virtual address** to determine which segment it belongs to?
  - A virtual address  $v'$  can be **divided** into two parts
  - **Top few bits** are for determining which segment this virtual address  $v'$  is in
  - Remaining bits are **the offset**, which points to the memory location **relative to** the segment base address

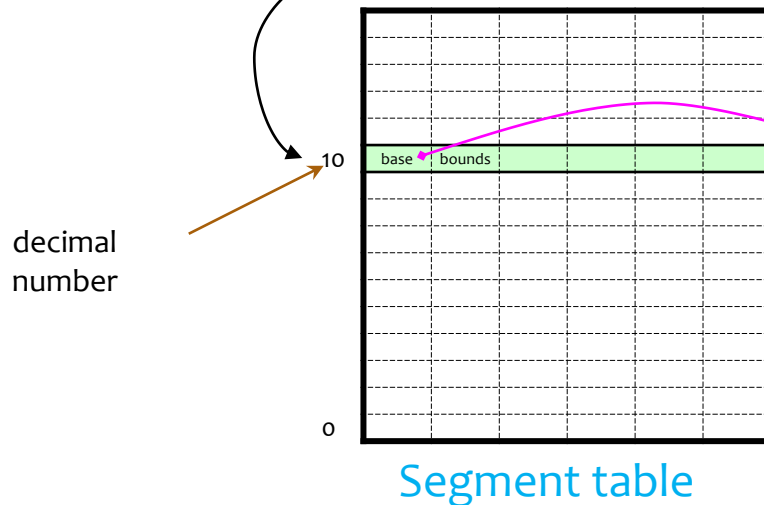
# Address Translation

Assume 16-bit virtual address 41414, the system uses top 4 bits to select the segment

No. of segments in  
a process  
 $= 2^4 = 16$



Max. size of each segment  
 $= 2^{12} = 4096$  bytes




+

physical address = base + offset

Each process has a segment table (may or may not be stored in MMU) that contains the address translation information of **all segments** in the process.




# Table Entry of Segment Table



- Protection bits
  - Indicate whether or not a program can **Read** or **Write** to a segment, or even **Execute** code in the segment
  - If process tries to write to a read-only segment, CPU would raise an exception to alert OS to deal with the offense
- Sharing of memory segment between processes
  - e.g., code sharing
  - By appropriate setting of protection bits, OS maps the **same** physical segment into **different** logical segments of **different** process address spaces.


# Table Entry of Segment Table



- **Grow bit**
  - Indicate which direction the segment grows
  - e.g., heap grows upwards, stack grows downwards
- **Valid bit**
  - When the segment is not in used, valid bit should set to false
  - If a process tries to access such memory, an exception will be generated

# External Fragmentation

- An issue arises from allocating memory to segments of **different sizes**
  - OS needs to find suitable space (block) **in physical memory** and allocates to new segments or to grow existing segments
  - Processes have multiple segments; when the process terminated, all physical memory used by its segments would be released
  - During normal runtime, many processes arrive, use up physical memory, terminate and release physical memory
    - So, where are those free slots? they appear in different places in the physical memory and are changing over time
  - Physical memory quickly becomes **full of little holes of free space**, which are **not large enough to hold a segment**
  - However, the sum of the holes might large enough to accommodate some more segments



# External Fragmentation

- Possible ways to combat external fragmentation
  - Coalescing
    - OS maintains a list of free memory blocks
    - When a process releases a segment, OS tries to see whether there are adjacent free blocks (i.e., they are in consecutive memory locations)
      - If so, combine adjacent free blocks into one large block
  - Compaction
    - Relocates all occupied areas of memory to one end
      - Forms a single contiguous block free space at the other end
    - Significant overhead
      - Processes have to be **suspended during relocation** and change their segment register values before resuming them
      - Extensive memory copy involved

# Managing Free Space

- OS needs to allocate physical memory to individual segment; thus, OS needs to know where to find free memory
- A straightforward approach is to use **a free-list**
  - Contains **references to all of the free chunks of space** in physical memory
  - OS just needs to search through the list to locate a suitable free block and allocates to a new request
- The Crux
  - What strategy can be used that is both **fast** and **minimizes external fragmentation**?

# Allocation Strategy – Best Fit

- Best-fit

- Search the free list and place the segment in the **smallest** possible memory block in which it will fit

- Example: **Freelist** → [16KB] → [5KB] → [30KB] → [14KB]

A new request for placing a segment of size 13KB

After allocation, the freelist becomes

**Freelist** → [16KB] → [5KB] → [30KB] → [1KB]

- By allocating a “just fit” block, we try to reduce wasted space
- However, still leave small amount of unused space
  - On long run, it too suffers with problem of numerous unusable small free areas
- Quite **significant performance overhead**
  - An **exhaustive search** for the best suit free block

# Allocation Strategy – Worst Fit

- Worst-fit
  - Place the segment in the **largest** available memory block in which it will fit
    - Example:  
**Freelist** → [16KB] → [5KB] → [30KB] → [14KB]  
  
After allocation, the freelist becomes  
**Freelist** → [16KB] → [5KB] → [17KB] → [14KB]
  - Leaves another “large” hole, making it more likely that another segment can fit in the hole
  - Quite **significant performance overhead**
    - A full search of the list is required

# Allocation Strategy – First Fit

- First-fit

- Place the segment in the **first memory** block (relative to the list head) in the free list in which it will fit

- Example:

**Freelist** → [16KB] → [5KB] → [30KB] → [14KB]

After allocation, the freelist becomes

**Freelist** → [3KB] → [5KB] → [30KB] → [14KB]

- Simple, **low execution-time overhead**
- This technique may create many small holes at the beginning of the free list



# Allocation Strategy – Next Fit

- Next-fit
  - Similar to first-fit
  - Instead, it searches the free list **at the point where the last search** was stopped at

- Example:           **Freelist** → [16KB] → **[5KB]** → [30KB] → [14KB]

After allocation, the freelist becomes

**Freelist** → [16KB] → [5KB] → **[17KB]** → [14KB]

- Somewhat surprisingly, best fit results in more wasted memory than **first fit and next fit** because it tends to fill up memory with tiny, useless holes. First fit generates larger holes on the average.

# Fixed-Size Blocks

- The only real solution to combat external fragmentation is to avoid the problem altogether, which is by **never allocating** physical memory in **variable-size blocks**
- Paging
  - The **whole process address space** is **divided into fixed-size memory block**, called (virtual) **page**
  - **Physical memory** is also divided into memory blocks of the same size of a virtual page – we call it a **page frame** or simply a **frame**

# Paging

- Advantages of Paging
  - Simply, because all pages/frames are of the same size
    - Completely **avoid external fragmentation** as each block is of the same size
      - Any virtual page of any process can be placed in any free page frame
  - This also **simplifies the free-space management**

# Address Translation – Page Table


- Also known as page map table
- **Each process** has its page table
  - Contains the address translation information of **every virtual page** of a process
- Let's have the calculation
  - With **32-bit** x86 processor, a process address space is of size  $2^{32}$  bytes = 4294967296 bytes = 4 **GiB**
  - The **default** size of a virtual page (as well as page frame) is  $2^{12}$  bytes = 4096 bytes = 4 **KiB**
  - How many virtual pages are there in one 32-bit address space?
    - $2^{32}/2^{12} = 2^{20}$  virtual pages
  - Thus, **a page table** needs to store  $2^{20} = 1048576$  entries of address translation information for a 32-bit process

# Address Translation


- Given a virtual address **v**, how to look up the translation information in the process's page table?
  - We need to **find out which virtual page** this virtual address is in
- Address Translation
  - Given a **virtual address v**, **split** it into two parts
  - **The top left p bits**, we call it the **virtual page number**, are for determining which virtual page this virtual address is in
  - Remaining bits on the right are **the offset d**, which points to the memory location relative to the starting address of a page/frame
  - How many virtual pages are there? And what is the size of a page?
    - No. of virtual pages =  $2^p$       Size of a page =  $2^d$  bytes

# Address Translation


- Assume 16-bit addressing and the page size is 1024 bytes =  $2^{10}$ 
  - No. of virtual pages =  $2^{(16-10)} = 64$  pages
- Given a virtual address 41414, we can determine its virtual page number  $p$  and the offset  $d$  in this virtual page



# Address Translation




# Page Table Entry (PTE)



- **Valid bit**
  - Indicates whether this PTE contains valid translation information
    - e.g., if this virtual page is unused, valid bit is set to false
  - Accessing an invalid page will generate an exception
- **Present bit**
  - Indicates whether this virtual page is in physical memory or on disk
  - If present bit is false, the value stored in PFN may not be correct or may not refer to a valid page frame



# Page Table Entry (PTE)



- **Dirty bit**
  - Indicates whether this virtual **page has been modified** since it was brought into memory
- **Reference** (or accessed or use) bit
  - Indicates whether this virtual page has been **recently accessed** by the process
  - As a useful indicator for the **page replacement strategies**

# Where are page tables stored?

- Page table is quite large
  - The no. of entries = no. of virtual pages in the process address space
  - Need to store contiguously in memory
    - 32-bit with 4KiB page size
    - $2^{32}/2^{12} = 2^{20}$  PTEs
    - assume each PTE occupies 4 bytes, we need 4 MiB for a page table
- Each process has its own page table
  - There are many processes running in the system !!!
- Must keep in physical memory – relatively slow speed of access

# Where are page tables stored?

- MMU needs to know **where to find the page table** for **current running process**
- CPU uses a **register** called **Page-table base register** to point to the **starting physical address** of the page table
  - OS must **reload this register during context switch** for the next process
- To translate each virtual address
  - MMU needs to **access page table** (via page-table base register) which is in the **physical memory**
  - Based on the virtual page number, reads the **PTE** and performs the translation
  - Fetches the desired data of that virtual address from the **physical memory**
  - In summary, **we need to access physical memory twice** for each memory reference

# The Crux

- How to make paging faster and page table smaller?
  - Each memory access results in two physical memory accesses
  - Page table must be stored contiguously in physical memory

# Summary

- Segmentation allows better utilization of physical memory as **only segments in use** are allocated with physical memory
- To support segmentation, MMU has to include more base-and-bounds register pairs (or even a segment table) for the address translation and protection
- Allocating variable-sized segments in physical memory leads to **external fragmentation**, which results in a waste of memory
- No matter how smart the allocation policies, external fragmentation still exists.

## Summary (2)

- By dividing the memory into **fixed-size memory blocks**, this avoids external fragmentation and simplifies the free-space management
- To translate a virtual address, MMU
  - Extracts the virtual page number from the virtual address
  - Accesses process's page table in physical memory to get the **physical page frame number**
  - Concatenates PFN with offset to form the physical address
  - Fetches the data from the physical memory